

```

1 # Steps to load extra files and download dataset.
2 import os
3 !git clone "https://vipulrawat008:Invincible007!@github.com/the-visionary/assignment1"
4 os.chdir('/content')
5 os.chdir('assignment1')
6 os.chdir('assignment1-part3')
7 assert (os.getcwd() == '/content/assignment1/assignment1-part3')
8
9 !bash download_data.sh
10 assert (os.getcwd() == '/content/assignment1/assignment1-part3')

```

```

Cloning into 'assignment1'...
remote: Enumerating objects: 59, done.
remote: Counting objects: 100% (59/59), done.
remote: Compressing objects: 100% (51/51), done.
remote: Total 59 (delta 20), reused 32 (delta 4), pack-reused 0
Unpacking objects: 100% (59/59), done.
--2021-03-06 03:57:23-- http://host.robots.ox.ac.uk/pascal/VOC/voc2007/VOCtrainval_06-Nov-2007.tar
Resolving host.robots.ox.ac.uk (host.robots.ox.ac.uk)... 129.67.94.152
Connecting to host.robots.ox.ac.uk (host.robots.ox.ac.uk)|129.67.94.152|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 460032000 (439M) [application/x-tar]
Saving to: 'VOCtrainval_06-Nov-2007.tar'

VOCtrainval_06-Nov- 100%[=====>] 438.72M  9.41MB/s   in 47s

2021-03-06 03:58:11 (9.26 MB/s) - 'VOCtrainval_06-Nov-2007.tar' saved [460032000/460032000]

--2021-03-06 03:58:13-- http://host.robots.ox.ac.uk/pascal/VOC/voc2007/VOCtest_06-Nov-2007.tar
Resolving host.robots.ox.ac.uk (host.robots.ox.ac.uk)... 129.67.94.152
Connecting to host.robots.ox.ac.uk (host.robots.ox.ac.uk)|129.67.94.152|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 451020800 (430M) [application/x-tar]
Saving to: 'VOCtest_06-Nov-2007.tar'

VOCtest_06-Nov-2007 100%[=====>] 430.13M  9.16MB/s   in 47s

2021-03-06 03:59:00 (9.10 MB/s) - 'VOCtest_06-Nov-2007.tar' saved [451020800/451020800]

```

```

1 import os
2 os.chdir('/content')
3 os.chdir('assignment1')
4 os.chdir('assignment1-part3')

```

Assignment 1 Part 3B: Developing Your Own Classifier

```

1 import os
2 import numpy as np
3 import torch
4 import torch.nn as nn
5 import torchvision
6

```

```

7 from torchvision import transforms
8 from sklearn.metrics import average_precision_score
9 from PIL import Image, ImageDraw
10 import matplotlib.pyplot as plt
11 from classifier import SimpleClassifier, Classifier#, AlexNet
12 from voc_dataloader import VocDataset, VOC_CLASSES
13
14 %matplotlib inline
15 %load_ext autoreload
16 # %reload_ext autoreload
17 %autoreload 2

```

Part 3B: Design your own network

In this notebook, your task is to create and train your own model for multi-label classification on VOC Pascal

What to do

1. You will make change on network architecture in `classifier.py`.
2. You may also want to change other hyperparameters to assist your training to get a better performance instructions.

What to submit

Check the submission template for details what to submit.

```

1 def train_classifier(train_loader, classifier, criterion, optimizer):
2     classifier.train()
3     loss_ = 0.0
4     losses = []
5     for i, (images, labels) in enumerate(train_loader):
6         images, labels = images.to(device), labels.to(device)
7         optimizer.zero_grad()
8         logits = classifier(images)
9         loss = criterion(logits, labels)
10        loss.backward()
11        optimizer.step()
12        losses.append(loss)
13    return torch.stack(losses).mean().item()

```

```

1 def test_classifier(test_loader, classifier, criterion, print_ind_classes=True, print_loss=True):
2     classifier.eval()
3     losses = []
4     with torch.no_grad():
5         y_true = np.zeros((0,21))
6         y_score = np.zeros((0,21))
7         for i, (images, labels) in enumerate(test_loader):
8             images, labels = images.to(device), labels.to(device)
9             logits = classifier(images)
10            y_true = np.concatenate((y_true, labels.cpu().numpy()), axis=0)

```

```

11         y_score = np.concatenate((y_score, logits.cpu().numpy()), axis=0)
12         loss = criterion(logits, labels)
13         losses.append(loss.item())
14         aps = []
15         # ignore first class which is background
16         for i in range(1, y_true.shape[1]):
17             ap = average_precision_score(y_true[:, i], y_score[:, i])
18             if print_ind_classes:
19                 print('----- Class: {:<12}      AP: {:>8.4f} -----'.format(VOC_
20                 aps.append(ap)
21
22         mAP = np.mean(aps)
23         test_loss = np.mean(losses)
24         if print_total:
25             print('mAP: {0:.4f}'.format(mAP))
26             print('Avg loss: {}'.format(test_loss))
27
28     return mAP, test_loss, aps

```

```

1 def plot_losses(train, val, test_frequency, num_epochs):
2     plt.plot(train, label="train")
3     indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0)]
4     plt.plot(indices, val, label="val")
5     plt.title("Loss Plot")
6     plt.ylabel("Loss")
7     plt.xlabel("Epoch")
8     plt.legend()
9     plt.show()
10
11 def plot_mAP(train, val, test_frequency, num_epochs):
12     indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0)]
13     plt.plot(indices, train, label="train")
14     plt.plot(indices, val, label="val")
15     plt.title("mAP Plot")
16     plt.ylabel("mAP")
17     plt.xlabel("Epoch")
18     plt.legend()
19     plt.show()
20

```

```

1
2 def train(classifier, num_epochs, train_loader, val_loader, criterion, optimizer, te
3     train_losses = []
4     train_mAPs = []
5     val_losses = []
6     val_mAPs = []
7
8     for epoch in range(1,num_epochs+1):
9         print("Starting epoch number " + str(epoch))
10        train_loss = train_classifier(train_loader, classifier, criterion, optimizer
11        train_losses.append(train_loss)
12        print("Loss for Training on Epoch " +str(epoch) + " is " + str(train_loss))
13        if(epoch%test_frequency==0 or epoch==1):

```

```

13         if (epoch % test_frequency == 0 or epoch == 1):
14             mAP_train, _, _ = test_classifier(train_loader, classifier, criterion, F
15             train_mAPs.append(mAP_train)
16             mAP_val, val_loss, _ = test_classifier(val_loader, classifier, criterion
17             print('Evaluating classifier')
18             print("Mean Precision Score for Testing on Epoch " + str(epoch) + " is " +
19             val_losses.append(val_loss)
20             val_mAPs.append(mAP_val)
21
22     return classifier, train_losses, val_losses, train_mAPs, val_mAPs

```

▼ Developing Your Own Model

▼ Goal

To meet the benchmark for this assignment you will need to improve the network. Note you should have not trained AlexNet really well, but training AlexNet from scratch performs much worse. We hope you can design a better architecture and AlexNet to train from scratch.

How to start

You may take inspiration from other published architectures and architectures discussed in lecture. However, you should not use predefined models (e.g. models from torchvision) or use pretrained weights. Training must be done from scratch.

Some hints

There are a variety of different approaches you should try to improve performance from the simple classifier.

- Network architecture changes
 - Number of layers: try adding layers to make your network deeper
 - Batch normalization: adding batch norm between layers will likely give you a significant performance improvement
 - Residual connections: as you increase the depth of your network, you will find that having residual connections in your architectures will be helpful
- Optimizer: Instead of plain SGD, you may want to add a learning rate schedule, add momentum, or use an optimizer you have learned about like Adam. Check the `torch.optim` package for other optimizers
- Data augmentation: You should use the `torchvision.transforms` module to try adding random resized crops to your input data. Check `transforms.RandomResizedCrop` and `transforms.RandomHorizontalFlip` for this. Feel free to experiment with other data augmentation which can lead to better performance.
- Epochs: Once you have found a generally good hyperparameter setting try training for more epochs
- Loss function: You might want to add weighting to the `MultiLabelSoftMarginLoss` for classes that are rare. Feel free to experiment with a different loss function

Note

We will soon be providing some initial expectations of mAP values as a function of epoch so you can get an idea of how well your implementation works without waiting a long time for training to converge.

What to submit

Submit your best model and save all plots for the writeup.

```
1 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
2
3 normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
4                                   std= [0.229, 0.224, 0.225])
5
6 train_transform = transforms.Compose([
7     transforms.Resize(227),
8     transforms.CenterCrop(227),
9     transforms.ToTensor(),
10    normalize
11    ])
12
13 test_transform = transforms.Compose([
14     transforms.Resize(227),
15     transforms.CenterCrop(227),
16     transforms.ToTensor(),
17     normalize,
18    ])
19
20 ds_train = VocDataset('VOCdevkit_2007/VOC2007/', 'train', train_transform)
21 ds_val = VocDataset('VOCdevkit_2007/VOC2007/', 'val', test_transform)
22 ds_test = VocDataset('VOCdevkit_2007/VOC2007test/', 'test', test_transform)
23
```

```
/content/assignment1/assignment1-part3/voc_dataloader.py:109: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a common outcome of np.array.fromiter() calls) is deprecated. Use np.ndarray(...) instead.
return np.array(names), np.array(labels).astype(np.float32), np.array(box_indices), label_order
```

```
1 num_epochs = 15
2 test_frequency = 5
3 batch_size = 64
4
5 train_loader = torch.utils.data.DataLoader(dataset=ds_train,
6                                             batch_size=batch_size,
7                                             shuffle=True,
8                                             num_workers=1)
9
10 val_loader = torch.utils.data.DataLoader(dataset=ds_val,
11                                          batch_size=batch_size,
12                                          shuffle=True,
13                                          num_workers=1)
14
15 test_loader = torch.utils.data.DataLoader(dataset=ds_test,
16                                           batch_size=batch_size,
17                                           shuffle=False,
18                                           num_workers=1)
```

```
1 import torch
2 import torch.nn as nn
3 from torch.autograd import Variable
4 import torch.nn.functional as F
5 from torch import optim
```

```

6 import numpy as np
7
8 NUM_CLASSES = 21
9
10 class ResBlock(nn.Module):
11     def __init__(self, in_channels, out_channels, stride=1):
12         """
13         Args:
14             in_channels (int): Number of input channels.
15             out_channels (int): Number of output channels.
16             stride (int): Controls the stride.
17         """
18         super(ResBlock, self).__init__()
19
20         self.skip = nn.Sequential()
21
22         if stride != 1 or in_channels != out_channels:
23             self.skip = nn.Sequential(
24                 nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=3, stride=stride, padding=1),
25                 nn.BatchNorm2d(out_channels))
26         else:
27             self.skip = None
28
29         self.block = nn.Sequential(
30             nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=3, stride=1, padding=1),
31             nn.BatchNorm2d(out_channels),
32             nn.ReLU(),
33             nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=3, stride=1, padding=1),
34             nn.BatchNorm2d(out_channels))
35
36     def forward(self, x):
37         identity = x
38         out = self.block(x)
39
40         if self.skip is not None:
41             identity = self.skip(x)
42
43         out += identity
44         out = F.relu(out)
45
46         return out
47
48 class Classifier(nn.Module):
49     def __init__(self):
50         super(Classifier, self).__init__()
51         self.conv1 = nn.Conv2d(3, 32, 3, stride=1)
52         self.conv1b = nn.Conv2d(32, 32, 3, stride=2)
53         # self.conv2a = nn.Conv2d(48, 32, 1)
54         # self.conv2b = nn.Conv2d(32, 32, 5, stride=1, padding=2)
55         # self.conv2c = nn.Conv2d(32, 256, 1)
56         self.conv3a = nn.Conv2d(32, 64, 1)
57         self.res = ResBlock(64, 64)
58         self.conv3e = nn.Conv2d(64, 256, 1)
59         self.pool = nn.MaxPool2d(3, 2)

```

```

60         self.fc1 = nn.Linear(2*2*256, 200)
61         self.fc2 = nn.Linear(200, 100)
62         self.fc3 = nn.Linear(100, NUM_CLASSES)
63
64     def forward(self, x):
65         #print(x.size())
66         x = self.conv1(x)
67         x = F.relu(x)
68         #print(x.size())
69         x = self.pool(x)
70         #print(x.size())
71         x = self.conv1b(x)
72         x = F.relu(x)
73         #print(x.size())
74         x = self.conv1b(x)
75         x = F.relu(x)
76         #print(x.size())
77         x = self.pool(x)
78         #print(x.size())
79         x = self.conv3a(x)
80         x = F.relu(x)
81         #print(x.size())
82         x = self.res(x)
83         x = self.res(x)
84         x = self.res(x)
85         x = F.relu(x)
86         #print(x.size())
87         x = self.pool(x)
88         #print(x.size())
89         x = self.conv3e(x)
90         x = F.relu(x)
91         #print(x.size())
92         x = self.pool(x)
93         #print(x.size())
94         x = x.view(x.size()[0], 2*2*256)
95         x = F.relu(self.fc1(x))
96         x = F.relu(self.fc2(x))
97         x = self.fc3(x)
98         return x
99

```

```

1 # TODO: Run your own classifier here
2 classifier = Classifier().to(device)
3
4 criterion = nn.MultiLabelSoftMarginLoss()
5 optimizer = torch.optim.SGD(classifier.parameters(), lr=0.005, momentum=0.9)
6
7 classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier, num_e
8

```

Loss for Training on Epoch 10 is 0.231260746717453

----- Class: aeroplane AP: 0.1011 -----

----- Class: bicycle AP: 0.0557 -----

----- Class: bird AP: 0.0556 -----

-----	Class: boat	AP:	0.0596	-----
-----	Class: bottle	AP:	0.0516	-----
-----	Class: bus	AP:	0.0402	-----
-----	Class: car	AP:	0.1610	-----
-----	Class: cat	AP:	0.0726	-----
-----	Class: chair	AP:	0.1156	-----
-----	Class: cow	AP:	0.0304	-----
-----	Class: diningtable	AP:	0.0539	-----
-----	Class: dog	AP:	0.0883	-----
-----	Class: horse	AP:	0.0684	-----
-----	Class: motorbike	AP:	0.0778	-----
-----	Class: person	AP:	0.4932	-----
-----	Class: pottedplant	AP:	0.0521	-----
-----	Class: sheep	AP:	0.0162	-----
-----	Class: sofa	AP:	0.0662	-----
-----	Class: train	AP:	0.0723	-----
-----	Class: tvmonitor	AP:	0.0472	-----

mAP: 0.0890

Avg loss: 0.32551265358924864

Evaluating classifier

Mean Precision Score for Testing on Epoch 10 is 0.08895318486778567

Starting epoch number 11

Loss for Training on Epoch 11 is 0.23047009110450745

Starting epoch number 12

Loss for Training on Epoch 12 is 0.2282872200012207

Starting epoch number 13

Loss for Training on Epoch 13 is 0.22733020782470703

Starting epoch number 14

Loss for Training on Epoch 14 is 0.22457727789878845

Starting epoch number 15

Loss for Training on Epoch 15 is 0.2240038812160492

-----	Class: aeroplane	AP:	0.1863	-----
-----	Class: bicycle	AP:	0.0631	-----
-----	Class: bird	AP:	0.0555	-----
-----	Class: boat	AP:	0.1307	-----
-----	Class: bottle	AP:	0.0493	-----
-----	Class: bus	AP:	0.0588	-----
-----	Class: car	AP:	0.2048	-----
-----	Class: cat	AP:	0.0671	-----
-----	Class: chair	AP:	0.1170	-----
-----	Class: cow	AP:	0.0281	-----
-----	Class: diningtable	AP:	0.0494	-----
-----	Class: dog	AP:	0.0760	-----
-----	Class: horse	AP:	0.0763	-----
-----	Class: motorbike	AP:	0.0940	-----
-----	Class: person	AP:	0.5150	-----
-----	Class: pottedplant	AP:	0.0475	-----
-----	Class: sheep	AP:	0.0143	-----
-----	Class: sofa	AP:	0.0600	-----
-----	Class: train	AP:	0.0920	-----
-----	Class: tvmonitor	AP:	0.0469	-----

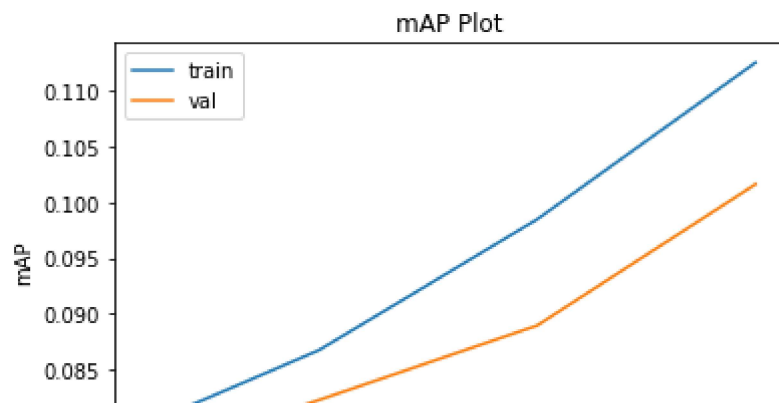

mAP: 0.1016

Avg loss: 0.3349403478205204

Evaluating classifier

Mean Precision Score for Testing on Epoch 15 is 0.10161171225786791

```
1 plot_losses(train_losses, val_losses, test_frequency, num_epochs)
2 plot_mAP(train_mAPs, val_mAPs, test_frequency, num_epochs)
```

```
1 mAP_test, test_loss, test_aps = test_classifier(test_loader, classifier, criterion)
2 print(mAP_test)
```

```
----- Class: aeroplane      AP:  0.2006 -----
----- Class: bicycle       AP:  0.0592 -----
----- Class: bird          AP:  0.0556 -----
----- Class: boat          AP:  0.1014 -----
----- Class: bottle        AP:  0.0618 -----
----- Class: bus           AP:  0.0513 -----
----- Class: car           AP:  0.2233 -----
----- Class: cat           AP:  0.0587 -----
----- Class: chair         AP:  0.1146 -----
----- Class: cow           AP:  0.0237 -----
----- Class: diningtable   AP:  0.0592 -----
----- Class: dog           AP:  0.0713 -----
----- Class: horse         AP:  0.0682 -----
----- Class: motorbike     AP:  0.0846 -----
----- Class: person        AP:  0.5183 -----
----- Class: pottedplant    AP:  0.0542 -----
----- Class: sheep         AP:  0.0133 -----
----- Class: sofa          AP:  0.0593 -----
----- Class: train         AP:  0.0883 -----
----- Class: tvmonitor     AP:  0.0608 -----
mAP: 0.1014
Avg loss: 0.33337560563515395
0.10139077554847706
```

```
1 torch.save(classifier.state_dict(), './voc_my_best_classifier.pth')
2
3 # output_submission_csv('my_solution.csv', test_aps)
```

