


Convolutional neural networks




Outline

- Building blocks
 - Convolutional layers and backprop rules
 - Pooling layers and nonlinearities
- Architectures:
 - 2012: AlexNet
 - 2013: ZFNet
 - 2014: VGGNet, GoogLeNet
 - 2015: ResNet
 - 2016: ResNeXt, DenseNet
 - etc.

This Class


Neural Network and Image

- Dimensionality
- Local relationship




Convolutional Neural Network (CNN)

- Convolution Layer
- Non-linearity Layer
- Pooling Layer
- Fully Connected Layer
- Classification Layer





ImageNet Challenge


- Progress
- Human Level Performance



Neural Networks



Multi-layer Neural Network & Image




How to apply NN over Image?

Multi-layer Neural Network & Image




Stretch pixels in
single column
vector



Multi-layer Neural Network & Image



Stretch pixels in
single column
vector




Problems ?

Multi-layer Neural Network & Image



Stretch pixels in
single column
vector

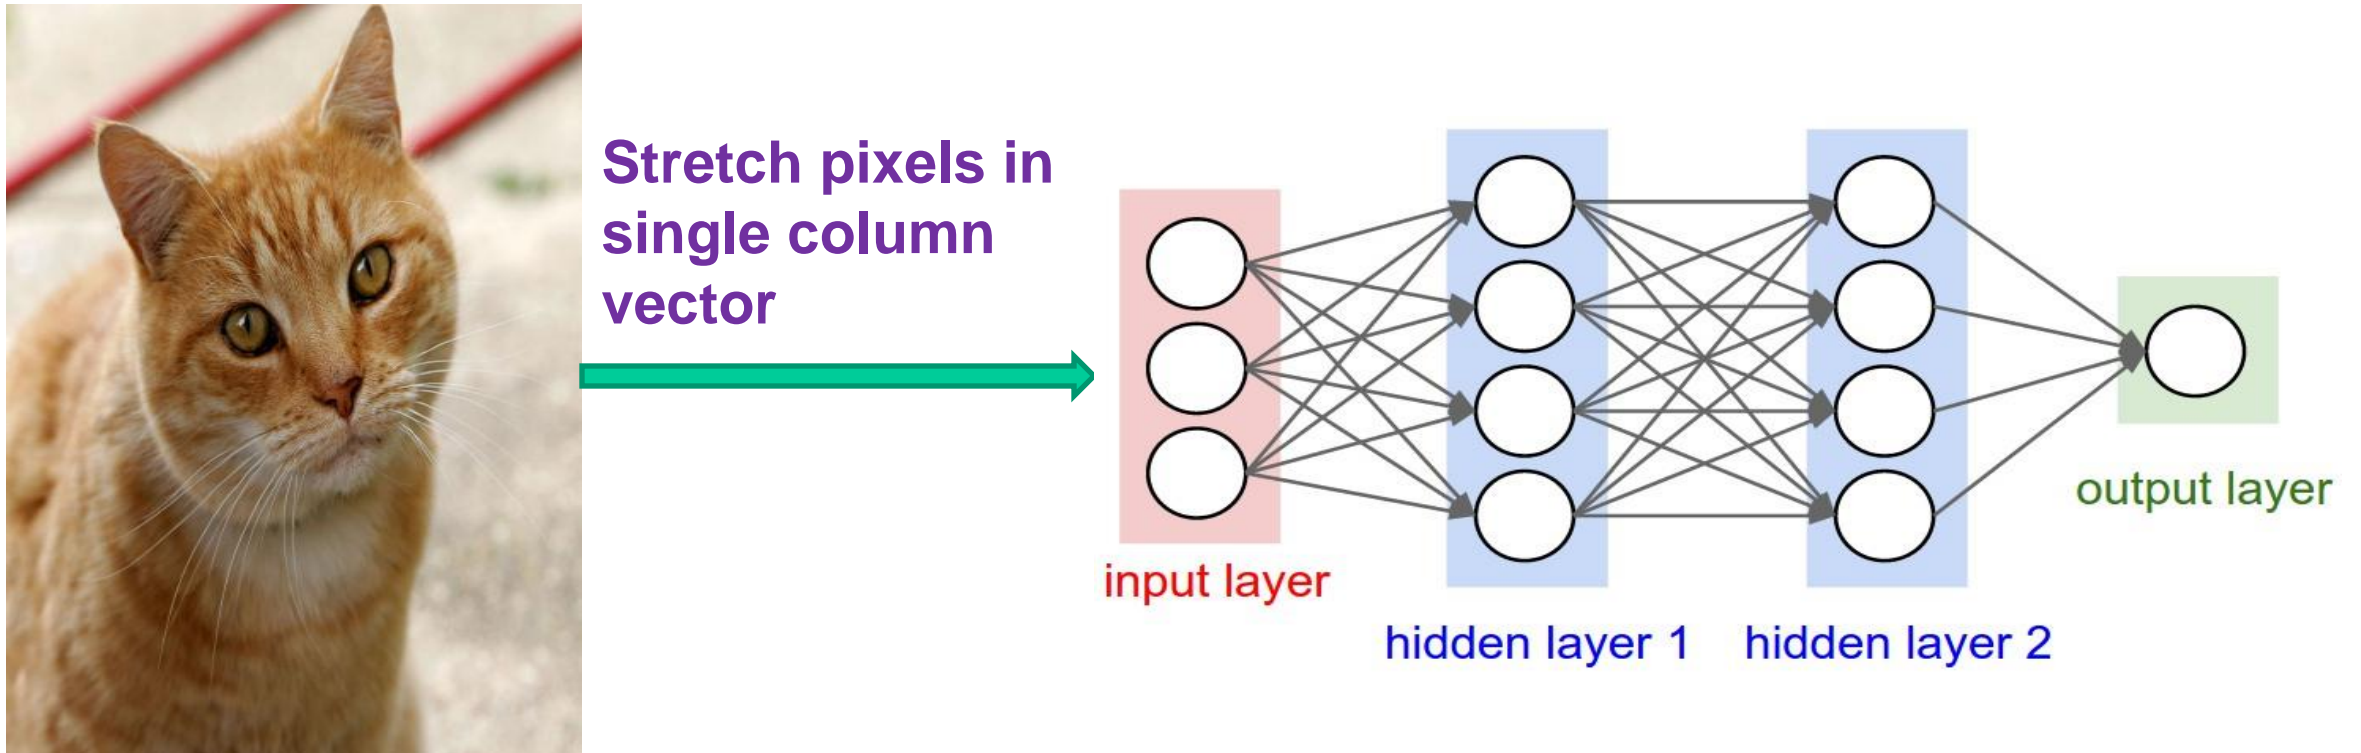


Problems:

High dimensionality

Local relationship

Multi-layer Neural Network & Image



Problems:

High dimensionality


Local relationship

Solution ?

Multi-layer Neural Network & Image



Stretch pixels in
single column
vector



Problems:

High dimensionality

Local relationship

Solution:

Convolutional Neural Network

Convolutional Neural Networks

Also known as

CNN,


ConvNet,

DCN

CNN = a multi-layer neural network with

1. Local connectivity
2. Weight sharing

CNN: Local Connectivity




Global connectivity

input units (neurons): 7

hidden units: 3


Hidden layer

Input layer



Local connectivity

CNN: Local Connectivity



Hidden layer

Input layer


Global connectivity

input units (neurons): 7

hidden units: 3


Number of parameters

- Global connectivity: ?
- Local connectivity: ?



Local connectivity

CNN: Local Connectivity



Hidden layer

Input layer


Global connectivity

input units (neurons): 7

hidden units: 3


Number of parameters

- Global connectivity: $3 \times 7 = 21$
- Local connectivity: $3 \times 3 = 9$



Local connectivity

CNN: Weight Sharing




Without weight sharing

- # input units (neurons): 7
- # hidden units: 3


Hidden layer

Input layer




With weight sharing

CNN: Weight Sharing




Without weight sharing




With weight sharing

- # input units (neurons): 7
- # hidden units: 3
- **Number of parameters**
 - Without weight sharing: ?
 - With weight sharing : ?

CNN: Weight Sharing




Without weight sharing



With weight sharing

- # input units (neurons): 7
- # hidden units: 3
- **Number of parameters**
 - Without weight sharing: $3 \times 3 = 9$
 - With weight sharing : $3 \times 1 = 3$

Convolutional Neural Networks



Layers used to build ConvNets

Input Layer (Input image)

Convolutional Layer

Non-linearity Layer (such as Sigmoid, Tanh, ReLU, PReLU, ELU, Swish, etc.)

Pooling Layer (such as Max Pooling, Average Pooling, etc.)


Fully-Connected Layer

Classification Layer (Softmax, etc.)

Convolutional Layer


32×32×3 Image

-> preserve spatial structure




Convolutional Layer

32×32×3 Image



5×5×3 Filter




Convolve the filter with the image i.e. “slide over the image spatially, computing dot products”

Convolutional Layer


Handling multiple input channels

Filters always extend the full depth of the input volume

32×32×3 Image




5×5×3 Filter



Convolve the filter with the image i.e. “slide over the image spatially, computing dot products”


Convolutional Layer

32×32×3 Image



Convolutional Layer

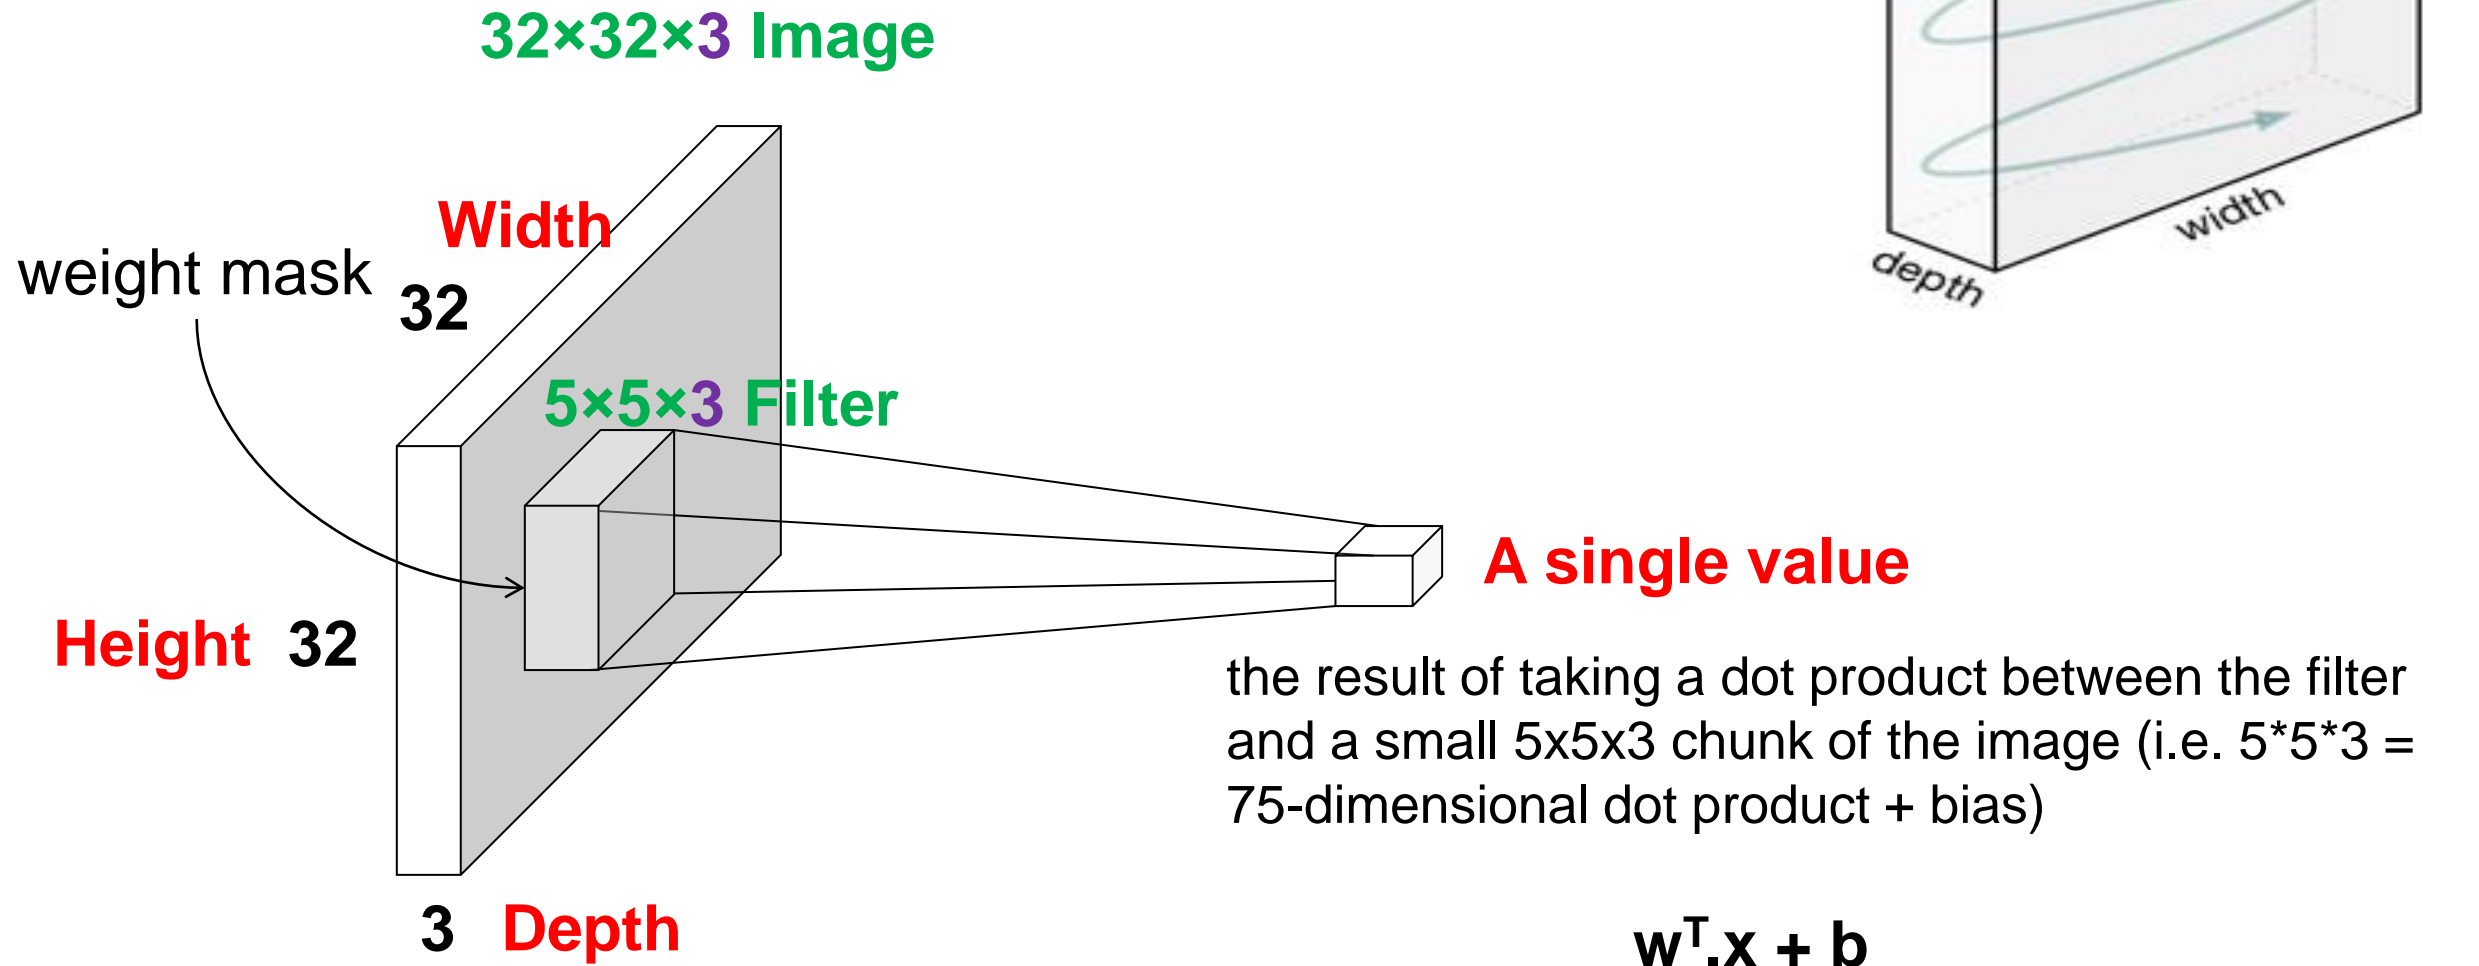
32×32×3 Image




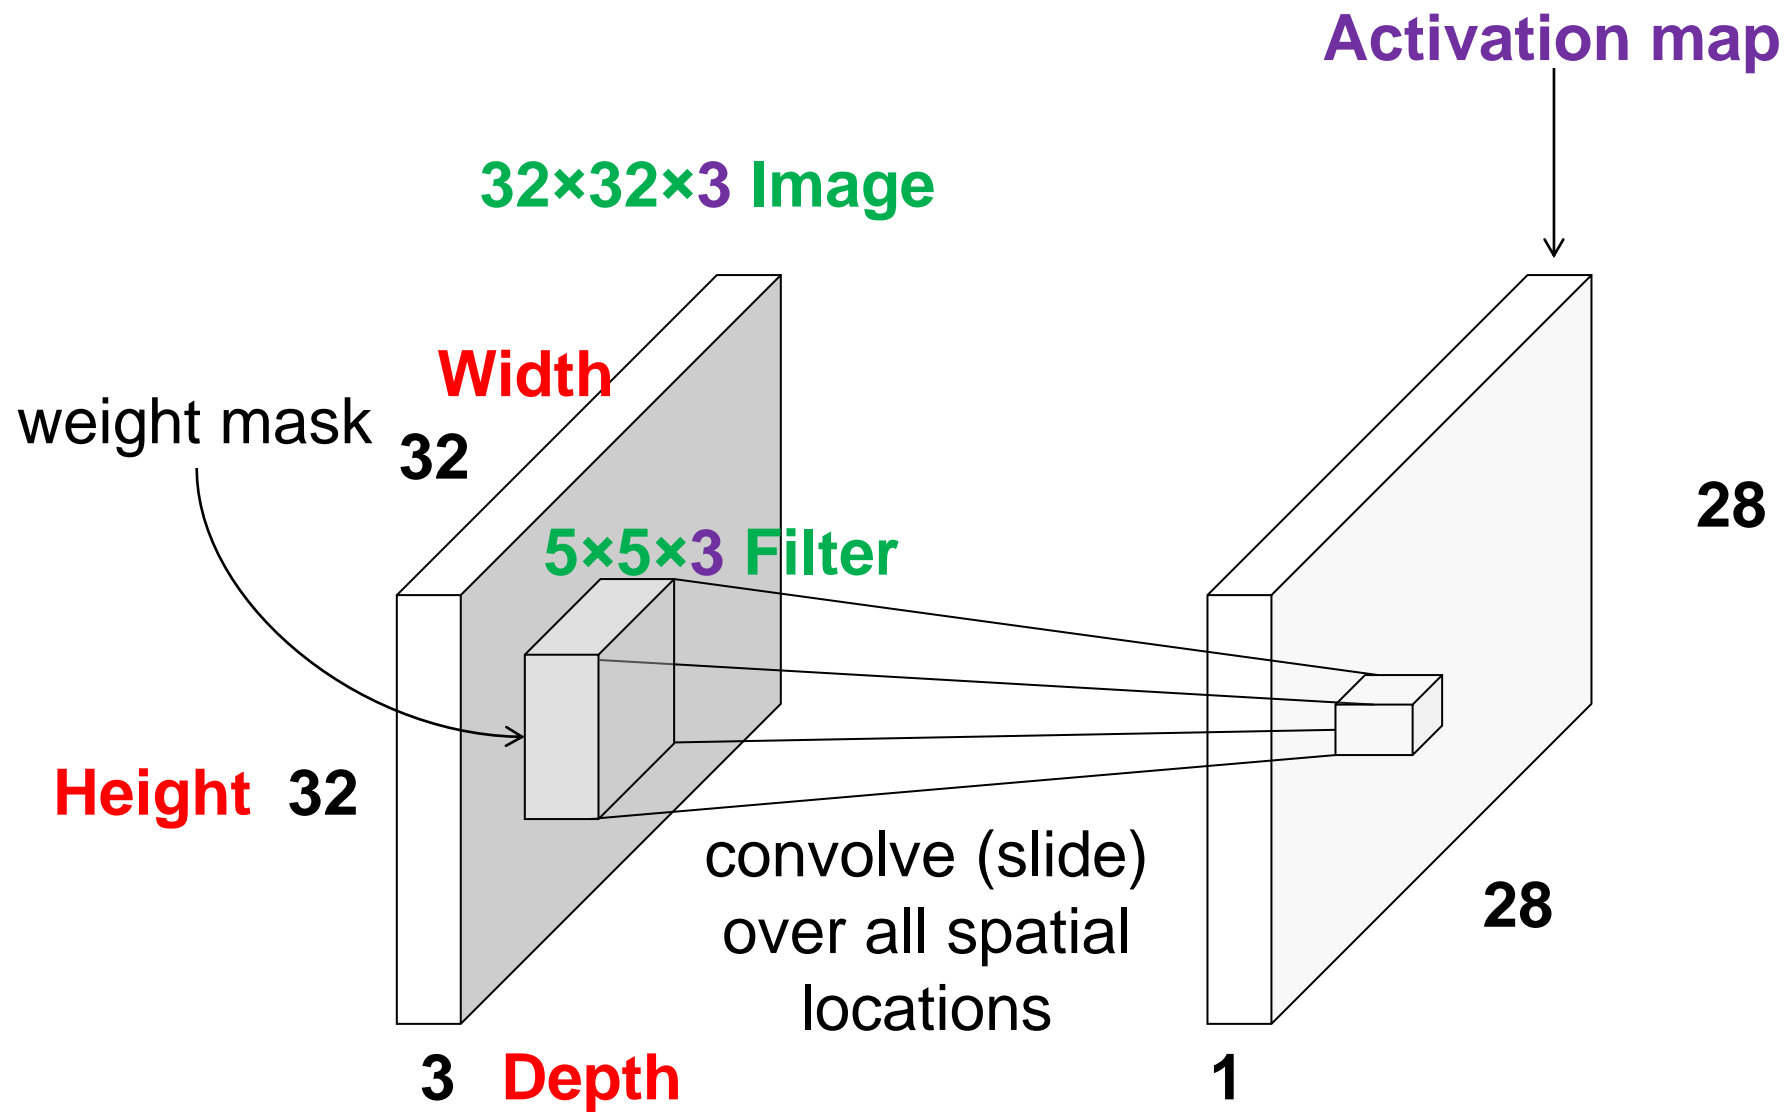
the result of taking a dot product between the filter and a small 5x5x3 chunk of the image (i.e. 5*5*3 = 75-dimensional dot product + bias)

$$w^T \cdot x + b$$

Convolutional Layer



Convolutional Layer




Convolutional Layer

Handling multiple output maps

32×32×3 Image

Activation maps




Convolutional Layer

Handling multiple output maps

32×32×3 Image

Activation maps




Convolutional Layer

Handling multiple output maps


32×32×3 Image

Activation maps



Convolution and traditional feature extraction


bank of K filters



K feature maps



image



feature map

Input Volume (+pad 1) (7x7x3)

 $x[:, :, 0]$

0	0	0	0	0	0	0
0	0	0	2	0	2	0
0	1	2	2	2	0	0

0	0	2	1	0	0	0
0	0	2	0	2	1	0
0	1	0	2	2	0	0
0	0	0	0	0	0	0

 $x[:, :, 1]$

0	0	0	0	0	0	0
0	1	0	0	2	0	0
0	1	0	1	0	1	0

0	0	0	2	1	2	0
0	1	2	1	0	0	0
0	1	2	2	0	2	0
0	0	0	0	0	0	0

 $x[:, :, 2]$

0	0	0	0	0	0	0
0	0	2	2	1	0	0
0	0	2	0	0	1	0

0	1	1	1	2	1	0
0	2	0	1	2	2	0
0	2	1	1	0	0	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

 $w0[:, :, 0]$

-1	0	-1
1	-1	1
-1	0	-1

 $w0[:, :, 1]$

0	-1	1
1	1	1
-1	1	-1

 $w0[:, :, 2]$

0	0	0
-1	1	-1
1	1	-1

Bias b0 (1x1x1)

 $b0[:, :, 0]$

1

Filter W1 (3x3x3)

 $w1[:, :, 0]$

0	1	1
-1	-1	-1
1	-1	1

 $w1[:, :, 1]$

-1	0	-1
1	1	1
-1	1	-1

 $w1[:, :, 2]$

0	1	0
0	0	-1
1	-1	0

Bias b1 (1x1x1)

 $b1[:, :, 0]$

0

Output Volume (3x3x2)

 $o[:, :, 0]$

-3	-1	0
-1	-8	2
3	0	3

 $o[:, :, 1]$

1	4	2
-1	4	5
3	1	3

toggle movement

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$						
0	0	0	0	0	0	0
0	0	0	2	0	2	0
0	1	2	2	2	0	0
0	0	2	1	0	0	0
0	0	2	0	2	1	0
0	1	0	2	2	0	0
0	0	0	0	0	0	0
$x[:, :, 1]$						
0	0	0	0	0	0	0
0	1	0	0	2	0	0
0	1	0	1	0	1	0
0	0	0	2	1	2	0
0	1	2	1	0	0	0
0	1	2	2	0	2	0
0	0	0	0	0	0	0
$x[:, :, 2]$						
0	0	0	0	0	0	0
0	0	2	2	1	0	0
0	0	2	0	0	1	0
0	1	1	1	2	1	0
0	2	0	1	2	2	0
0	2	1	1	0	0	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$		
-1	0	-1
1	-1	1
-1	0	-1
$w0[:, :, 1]$		
0	-1	1
1	1	1
-1	1	-1
$w0[:, :, 2]$		
0	0	0
-1	1	-1
1	1	-1
Bias $b0$ (1x1x1)		
$b0[:, :, 0]$		
1		

Filter W1 (3x3x3)

$w1[:, :, 0]$		
0	1	1
-1	-1	-1
1	-1	1
$w1[:, :, 1]$		
-1	0	-1
1	1	1
-1	1	-1
$w1[:, :, 2]$		
0	1	0
0	0	-1
1	-1	0
Bias $b1$ (1x1x1)		
$b1[:, :, 0]$		
0		

Output Volume (3x3x2)

$o[:, :, 0]$		
-3	-1	0
-1	-8	2
3	0	3
$o[:, :, 1]$		
1	4	2
-1	4	5
3	1	3

toggle movement

Input Volume (+pad 1) (7x7x3)

 $x[:, :, 0]$

0	0	0	0	0	0	0
0	0	0	2	0	2	0
0	1	2	2	2	0	0
0	0	2	1	0	0	0
0	0	2	0	2	1	0
0	1	0	2	2	0	0
0	0	0	0	0	0	0

 $x[:, :, 1]$

0	0	0	0	0	0	0
0	1	0	0	2	0	0
0	1	0	1	0	1	0
0	0	0	2	1	2	0
0	1	2	1	0	0	0
0	1	2	2	0	2	0
0	0	0	0	0	0	0

 $x[:, :, 2]$

0	0	0	0	0	0	0
0	0	2	2	1	0	0
0	0	2	0	0	1	0
0	1	1	1	2	1	0
0	2	0	1	2	2	0
0	2	1	1	0	0	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

 $w0[:, :, 0]$

-1	0	-1
1	-1	1
-1	0	-1

 $w0[:, :, 1]$

0	-1	1
1	1	1
-1	1	-1

 $w0[:, :, 2]$

0	0	0
-1	1	-1
1	1	-1

Bias b0 (1x1x1)

 $b0[:, :, 0]$

1

Filter W1 (3x3x3)

 $w1[:, :, 0]$

0	1	1
-1	-1	-1
1	-1	1

 $w1[:, :, 1]$

-1	0	-1
1	1	1
-1	1	-1

 $w1[:, :, 2]$

0	1	0
0	0	-1
1	-1	0

Bias b1 (1x1x1)

 $b1[:, :, 0]$

0

Output Volume (3x3x2)

 $o[:, :, 0]$

-3	-1	0
-1	-8	2
3	0	3

 $o[:, :, 1]$

1	4	2
-1	4	5
3	1	3


toggle movement

Convolutional Layer

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions

32×32×3 Image

Activation maps




Convolutional Layer

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions

32×32×3 Image

Activation maps




Convolutional Layer

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions

32×32×3 Image

Activation maps

Deeper activation map




Convolutional Layer

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions

32×32×3 Image


Activation maps

Deeper activation maps



Multilayer Convolution

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



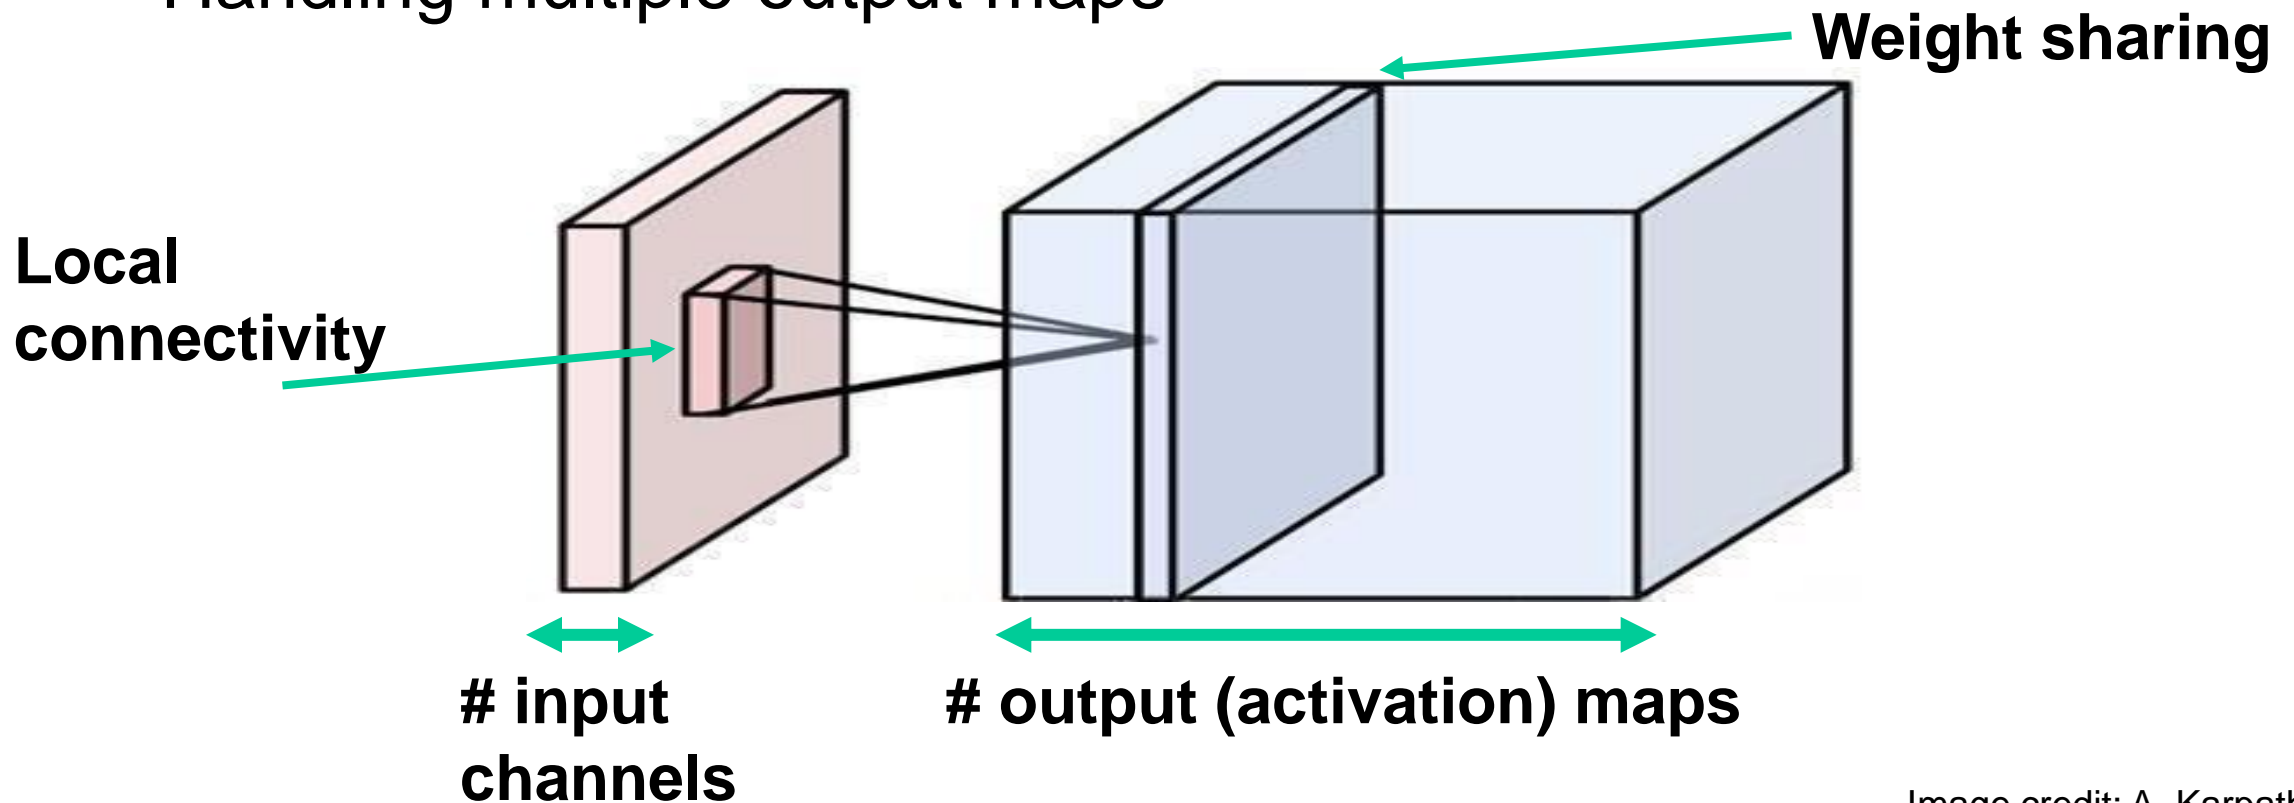
Any Convolution Layer

Local connectivity

Weight sharing

Handling multiple input channels

Handling multiple output maps



A closer look at spatial dimensions


7

7

7×7 input (spatially)
assume 3×3 filter

A closer look at spatial dimensions

7




7×7 input (spatially)
assume 3×3 filter

7

A closer look at spatial dimensions

7




7

7×7 input (spatially)
assume 3×3 filter

A closer look at spatial dimensions

7




7

7×7 input (spatially)
assume 3×3 filter

A closer look at spatial dimensions

7



7×7 input (spatially)
assume 3×3 filter

7



5×5 output

A closer look at spatial dimensions


7

7

7×7 input (spatially)
assume 3×3 filter
applied with stride 2

A closer look at spatial dimensions

7




7

7×7 input (spatially)
assume 3×3 filter
applied with stride 2

A closer look at spatial dimensions

7



7×7 input (spatially)
assume 3×3 filter
applied with stride 2

7



3×3 output

A closer look at spatial dimensions

7

7

7×7 input (spatially)
assume 3×3 filter
applied with stride 3

A closer look at spatial dimensions


7

7

7×7 input (spatially)
assume 3×3 filter
applied with stride 3

doesn't fit!
cannot apply 3x3 filter on
7x7 input with stride 3.


A closer look at spatial dimensions



Output size
 $(N - F) / \text{stride} + 1$

N e.g. $N = 7, F = 3$
stride 1 $\Rightarrow (7 - 3) / 1 + 1 = 5$
stride 2 $\Rightarrow (7 - 3) / 2 + 1 = 3$
stride 3 $\Rightarrow (7 - 3) / 3 + 1 = 2.33$

A closer look at spatial dimensions



E.g. 32x32 input convolved repeatedly with 5x5 filters shrinks volumes spatially! (32 -> 28 -> 24 ...). Shrinking too fast is not good, doesn't work well.

In practice: common to zero pad

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

e.g. input 7×7 (spatially)

3×3 filter, applied with stride 1

pad with 1 pixel border

What is the output dimension?

In practice: common to zero pad

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

e.g. input 7×7 (spatially)
3×3 filter, applied with stride 1
pad with 1 pixel border

7×7 Output

In practice: common to zero pad

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

e.g. input 7×7 (spatially)

3×3 filter, applied with stride 1

pad with 1 pixel border

7×7 Output

in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with

$(F-1)/2$. (will preserve size spatially)

e.g.

$F = 3 \Rightarrow$ zero pad with 1


$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Example

Input volume: 32x32x3
10 5x5 filters with stride 1, pad 2


Output volume size: ?



Example

Input volume: 32x32x3
10 5x5 filters with stride 1, pad 2

Output volume size:
 $(32+2*2-5)/1+1 = 32$ spatially,
so
32x32x10




Example

Input volume: 32x32x3

10 5x5 filters with stride 1, pad 2

Number of parameters in this layer?



Example

Input volume: 32x32x3


10 5x5 filters with stride 1, pad 2

Number of parameters in this layer?

each filter has

$5*5*3 + 1 = 76$ params (+1 for bias)

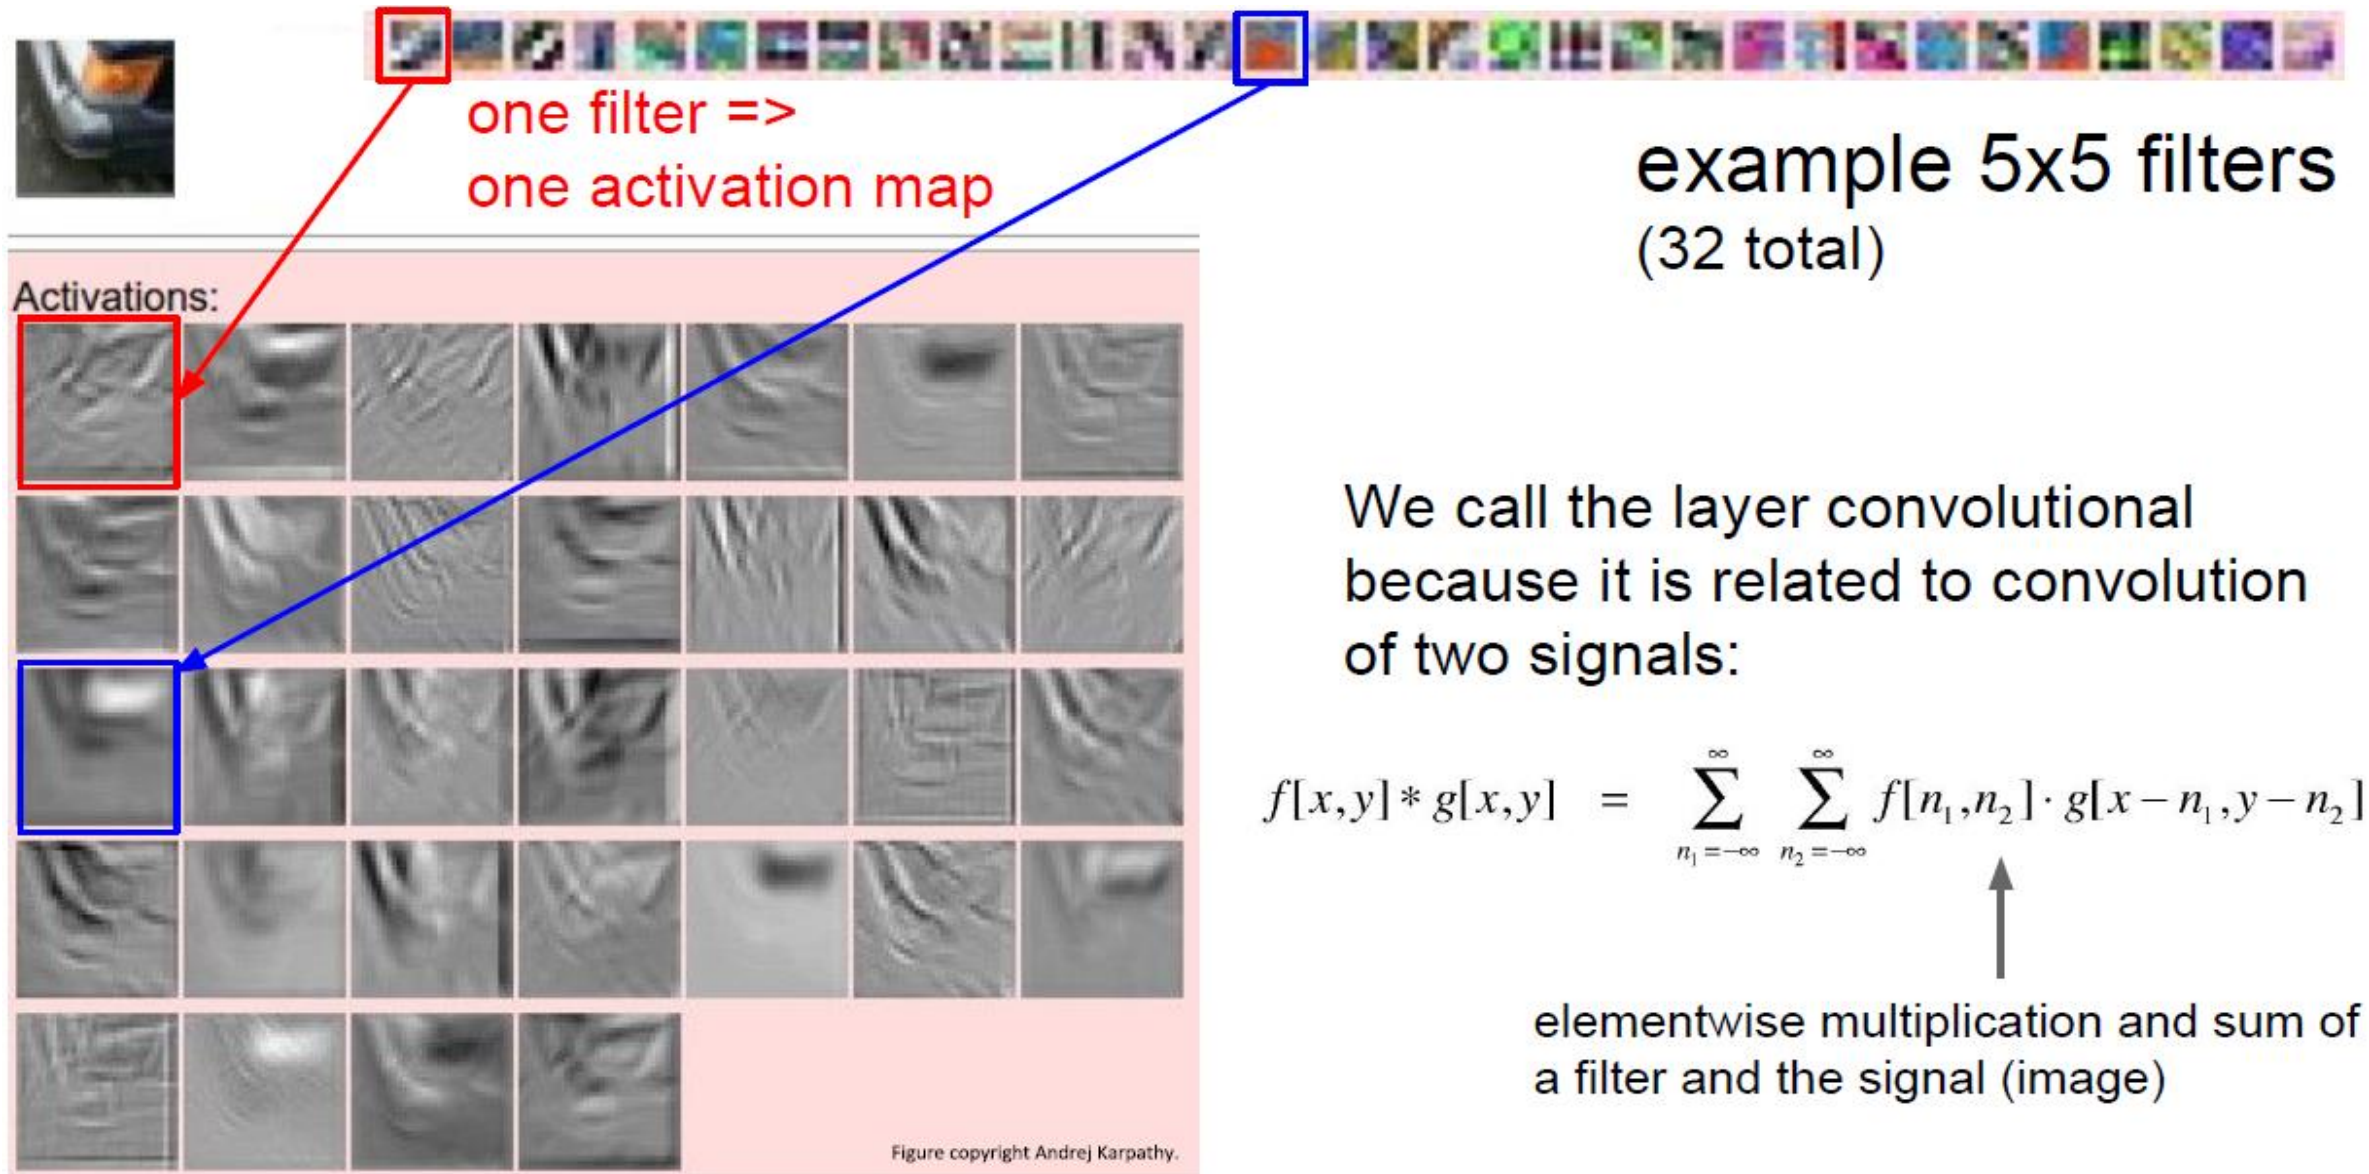
$\Rightarrow 76*10 = 760$



Summary. To summarize, the Conv Layer:

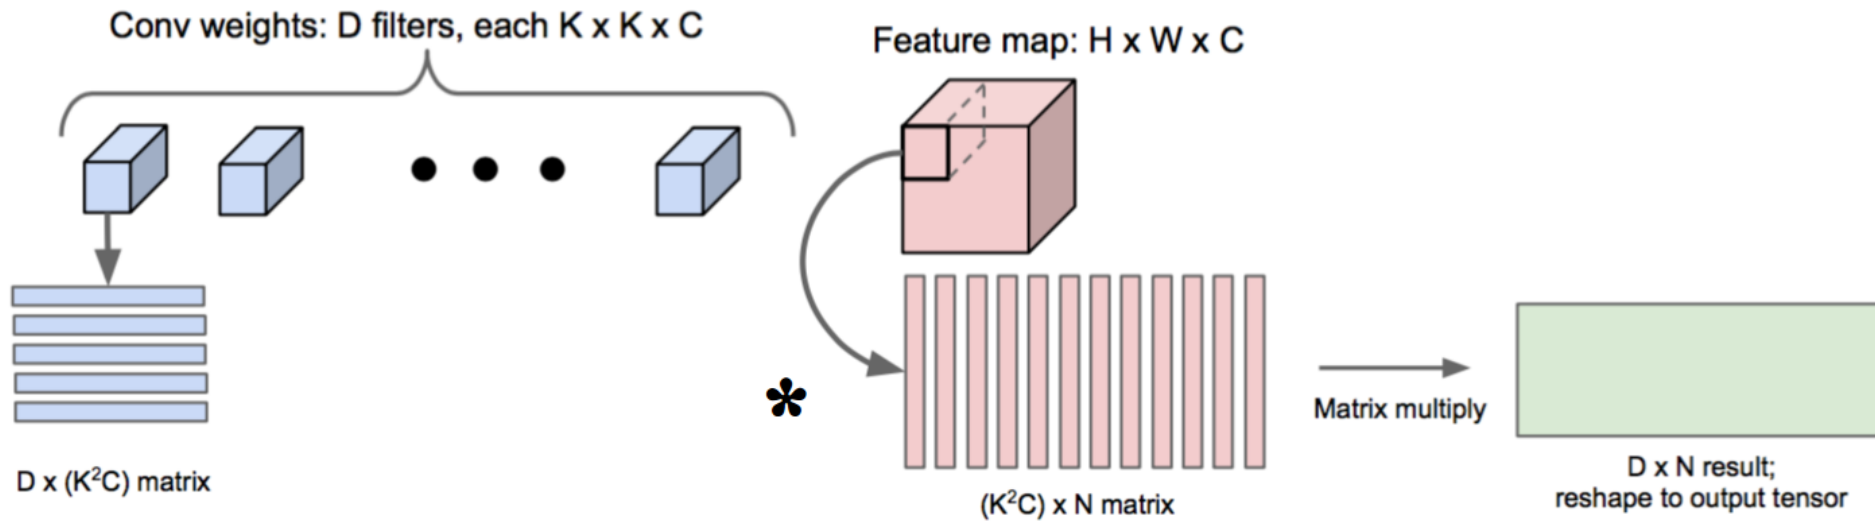
- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.

Convolution as feature extraction




Efficient implementation of convolutions

- Reshape all image neighborhoods into columns (im2col operation), do matrix-vector multiplication



Backpropagation for convolutional layer



Backpropagation for convolutional layer

$$\frac{\partial e}{\partial w_{ij}}$$


Backpropagation for convolutional layer

$$\frac{\partial e}{\partial w_{ij}} = \frac{\partial e}{\partial z} \frac{\partial z}{\partial w_{ij}} = \sum_{k,l} \frac{\partial e}{\partial z_{kl}} \frac{\partial z_{kl}}{\partial w_{ij}}$$

$$z_{kl} = \sum_{i,j=-f}^f w_{ij} x_{k+i, l+j}$$

For simplicity, assume filter indices go from $-f$ to f

$$\frac{\partial z_{kl}}{\partial w_{ij}} = x_{k+i, l+j}$$




Backpropagation for convolutional layer

$$\frac{\partial e}{\partial w_{ij}} = \frac{\partial e}{\partial z} \frac{\partial z}{\partial w_{ij}} = \sum_{k,l} \frac{\partial e}{\partial z_{kl}} \frac{\partial z_{kl}}{\partial w_{ij}} = \boxed{\sum_{k,l} \frac{\partial e}{\partial z_{kl}} x_{k+i, l+j}}$$

$$z_{kl} = \sum_{i,j=-f}^f w_{ij} x_{k+i, l+j}$$


For simplicity, assume filter indices go from $-f$ to f

$$\frac{\partial z_{kl}}{\partial w_{ij}} = x_{k+i, l+j}$$



Backpropagation for convolutional layer

$$\frac{\partial e}{\partial w_{ij}} = \sum_{k,l} \frac{\partial e}{\partial z_{kl}} x_{k+i, l+j}$$




Backpropagation for convolutional layer

$$\frac{\partial e}{\partial x_{kl}}$$


Backpropagation for convolutional layer

$$\frac{\partial e}{\partial x_{kl}} = \frac{\partial e}{\partial z} \frac{\partial z}{\partial x_{kl}} = \sum_{m,n} \frac{\partial e}{\partial z_{mn}} \frac{\partial z_{mn}}{\partial x_{kl}}$$



Backpropagation for convolutional layer


$$\frac{\partial e}{\partial x_{kl}} = \frac{\partial e}{\partial z} \frac{\partial z}{\partial x_{kl}} = \sum_{m,n} \frac{\partial e}{\partial z_{mn}} \frac{\partial z_{mn}}{\partial x_{kl}} = \sum_{i,j} \frac{\partial e}{\partial z_{k+i,l+j}} \frac{\partial z_{k+i,l+j}}{\partial x_{kl}}$$



Backpropagation for convolutional layer

$$\frac{\partial e}{\partial x_{kl}} = \sum_{i,j} \frac{\partial e}{\partial z_{k+i,l+j}} \frac{\partial z_{k+i,l+j}}{\partial x_{kl}}$$


$$\frac{\partial z_{k+i,l+j}}{\partial x_{kl}} = w_{-i,-j}$$



Backpropagation for convolutional layer

$$\frac{\partial e}{\partial x_{kl}} = \sum_{i,j} \frac{\partial e}{\partial z_{k+i,l+j}} \frac{\partial z_{k+i,l+j}}{\partial x_{kl}} = \boxed{\sum_{i,j} \frac{\partial e}{\partial z_{k+i,l+j}} w_{-i,-j}}$$

$$\frac{\partial z_{k+i,l+j}}{\partial x_{kl}} = w_{-i,-j}$$




Backpropagation for convolutional layer

$$\frac{\partial e}{\partial x_{kl}} = \sum_{i,j} \frac{\partial e}{\partial z_{k+i,l+j}} w_{-i,-j}$$


Backpropagation for convolutional layer

$$\frac{\partial e}{\partial x_{kl}} = \sum_{i,j} \frac{\partial e}{\partial z_{k+i,l+j}} w_{-i,-j}$$




Pooling Layer


- makes the representations smaller and more manageable
- operates over each activation map independently:



Max Pooling



max pool with 2x2 filters
and stride 2




Backward pass: upstream
gradient is passed back only to
the unit with max value

Pooling Layer

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires two hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

Fully Connected Layer


- **Connect every neuron in one layer to every neuron in another layer**
- **Same as the traditional multi-layer perceptron neural network**



Fully Connected Layer

- **Connect every neuron in one layer to every neuron in another layer**
- **Same as the traditional multi-layer perceptron neural network**

**No. of Neurons (Last FC)
= No. of classes**




Non-linearity Layer

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Classification/Loss Layer

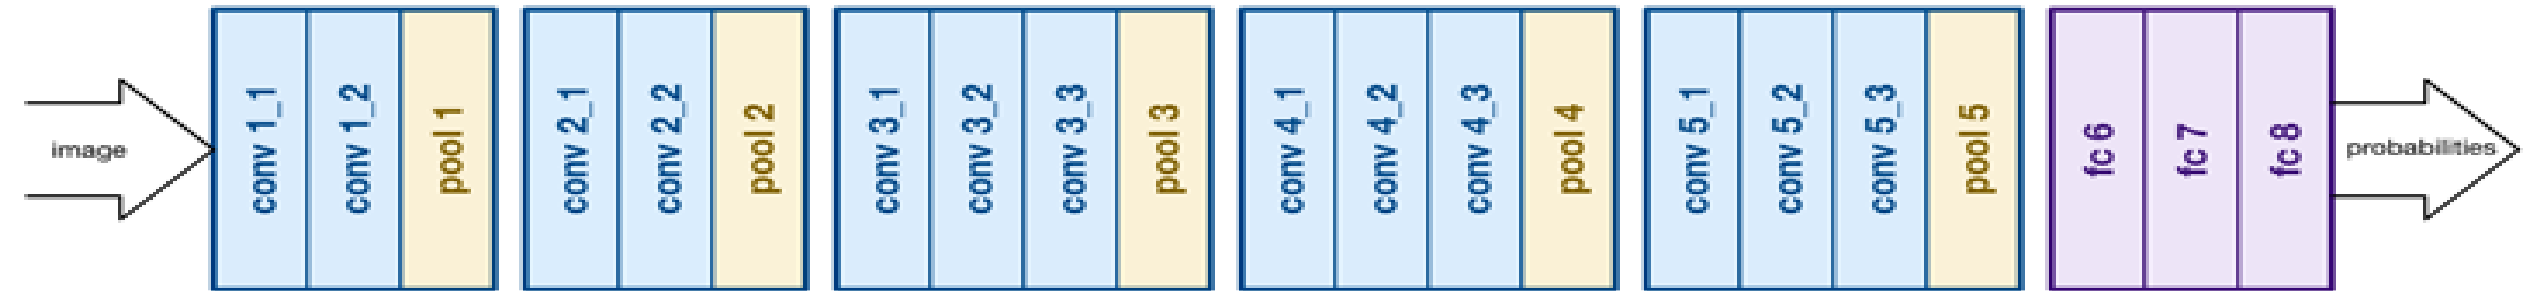
SVM Classifier

SVM Loss/Hinge Loss/Max-margin Loss

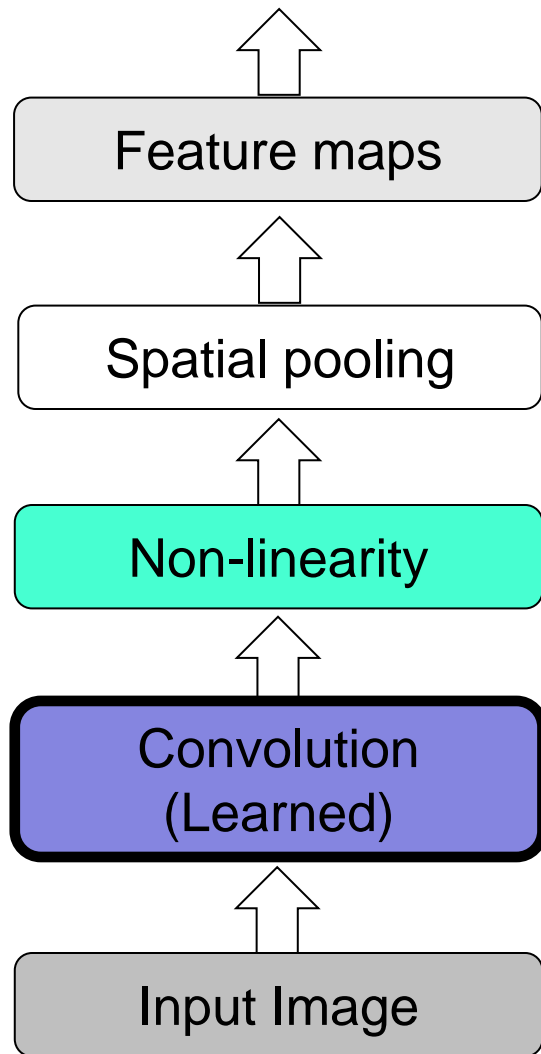
Softmax Classifier

Softmax Loss/Cross-entropy Loss

A typical CNN structure



Summary: CNN pipeline

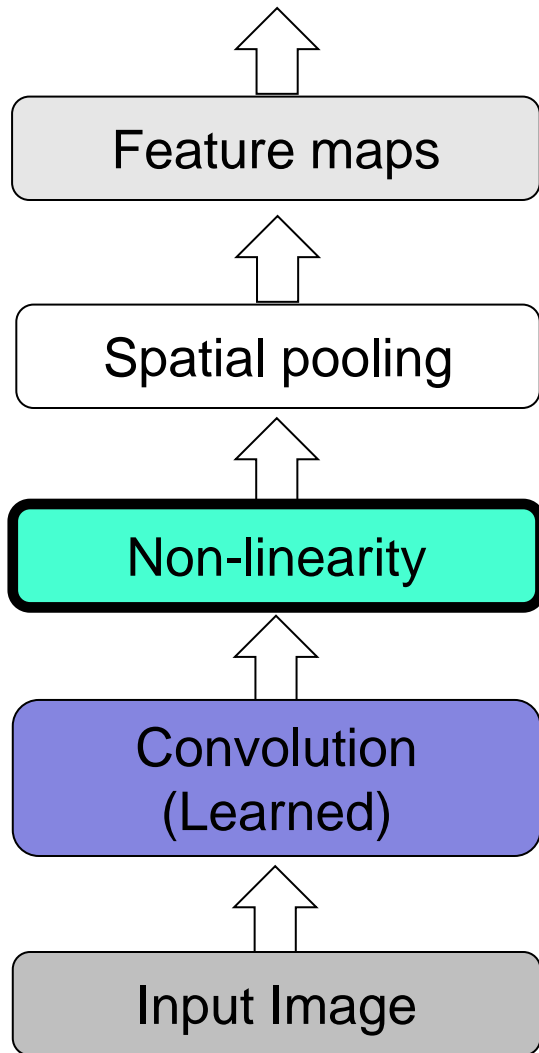


Input




Feature Map


Summary: CNN pipeline




ReLU
 $\max(0, x)$



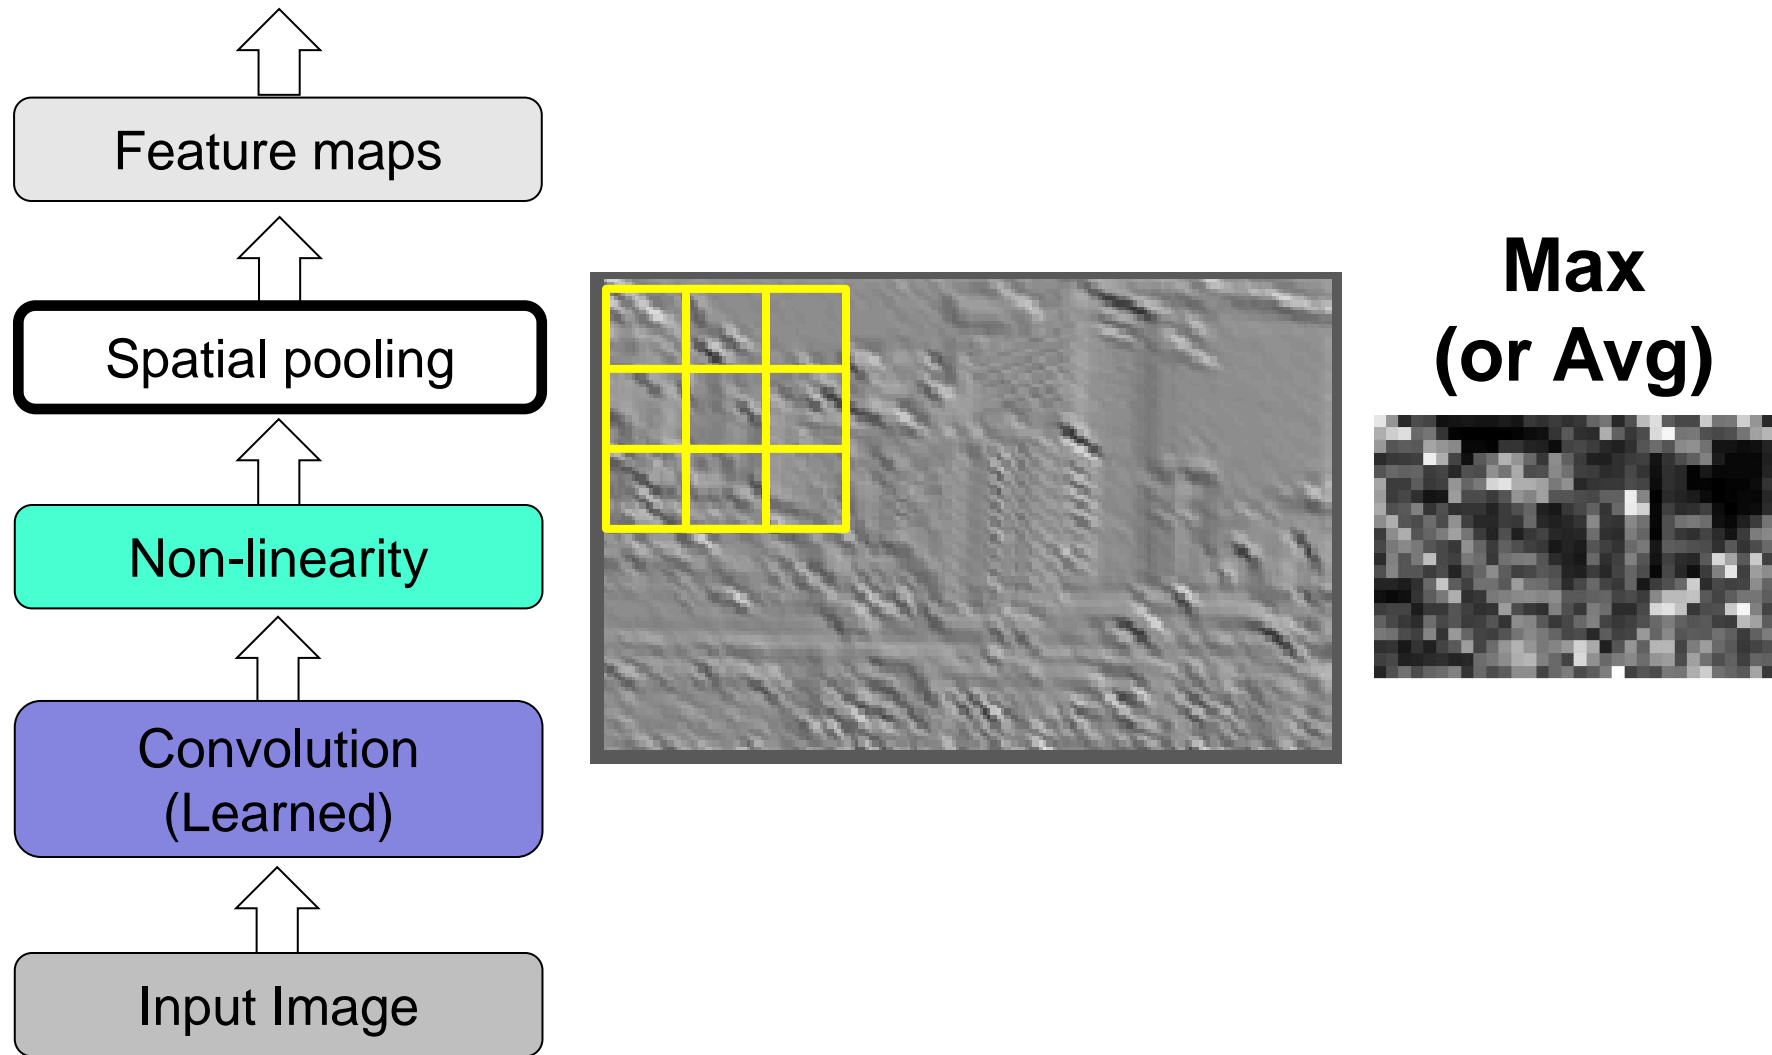
Leaky ReLU
 $\max(0.1x, x)$




ELU
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Summary: CNN pipeline




Summary: CNN pipeline



Softmax layer:

ImageNet Challenge

IMAGENET




- ~14 million labeled images, 20k classes
- Images gathered from Internet
- Human labels via Amazon MTurk
- Challenge: 1.2 million training images, 1000 classes

www.image-net.org/challenges/LSVRC/

ImageNet Challenge

IMAGENET




**GPUS
+
Data**

- 14 million labeled images, 20k classes
- Images gathered from Internet
- Human labels via Amazon MTurk
- Challenge: 1.2 million training images, 1000 classes

www.image-net.org/challenges/LSVRC/


Progress on ImageNet Challenge

ImageNet Image Classification Top5 Error



Progress on ImageNet Challenge

ImageNet Image Classification Top5 Error



Best Non-ConvNet in 2012: 26.2%

Things to remember

Neural network and Image

- Neuroscience, Perceptron, Problems due to High Dimensionality and Local Relationship

Convolutional neural network (CNN)

- Convolution Layer,
- Nonlinearity Layer,
- Pooling Layer,
- Fully Connected Layer,
- Loss/Classification Layer

Progress on ImageNet challenge

- Latest SENet, Winner 2017

Acknowledgement

Thanks to the following courses and corresponding researchers for making their teaching/research material online

- Deep Learning, Stanford University
- Introduction to Deep Learning, University of Illinois at Urbana-Champaign
- Introduction to Deep Learning, Carnegie Mellon University
- Convolutional Neural Networks for Visual Recognition, Stanford University
- Natural Language Processing with Deep Learning, Stanford University
- And Many More

Next Classes

Training Aspects of CNN

Activation Functions

Dataset Preparation

Data Preprocessing

Weight Initialization

Optimization Methods

Learning Rate

Transfer Learning

Generalization

