

Message Ordering and Group Communications

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

About this topic

This course covers various concepts in **Message Ordering and Group communication in Distributed Systems**. We will also focus on different models of communications and their pros and cons

What did you learn so far?

- Challenges in Message Passing systems
- Distributed Sorting
- Space-Time Diagram
- Partial Ordering / Total Ordering
- Causal Ordering - Precedence Relations
- Concurrent Events
- Local Clocks and Vector Clocks
- Distributed Snapshots
- Termination Detection using Dist. Snapshots
- Leader Election Problem in Rings

Recent Topic ...

→ Topology Abstraction and Overlays

- Various Interconnection Topologies
 - Abstraction - Basic Concepts
 - Interconnection Patterns suitable for message propagation
 - Types of Algorithms and their executions
 - Measures and Metrics
- Many more to come up ... stay tuned in !!

Topics to focus on ...

- Leader Election in Distributed Systems
- Topology Abstraction and Overlays
- Message Ordering
- Group Communication
- Distributed Mutual Exclusion
- Deadlock Detection
- Check pointing and rollback recovery

For End Semester

Message Ordering / Group Communication

Models of Communication

→ One - to - One

→ Unicast

→ 1 - 1

→ Point - to - point

→ Anycast

→ 1 - nearest 1 of several identical nodes

→ One - to - Many

→ Multicast

→ 1 - many

→ Group Communication

→ Broadcast

→ 1 - All

Groups

- Why groups?
 - Groups allow us to deal with a collection of processes as one abstraction
- Send message to one entity
 - Deliver to entire group
- Groups are dynamic
 - Created and destroyed
 - Processes can join or leave
 - May belong to 0 or more groups
- Primitives
 - `join_group`, `leave_group`, `send_to_group`, `query_membership`

Design Issues

Closed vs. Open

- Closed: only group members can sent messages

Peer vs. Hierarchical

- Peer: each member communicates with group
- Hierarchical: go through dedicated coordinator(s)
- Diffusion: send to other servers & clients

Managing membership & group creation/deletion

- Distributed vs. centralized

Leaving & joining must be synchronous

Fault tolerance


- Reliable message delivery? What about missing members?

Failures

- **Crash failure**
 - Process stops communicating
- **Omission failure (typically due to network)**
 - Send omission: A process fails to send messages
 - Receive omission: A process fails to receive messages
- **Byzantine Failure**
 - Some messages are faulty, including sending fake messages
- **Partition Failure**
 - The network may get segmented, dividing the group into two or more unreachable sub-groups


Multiple Unicasts

→ Sender knows Group members



Hierarchical

- Multiple unicasts via group coordinator
- coordinator knows group members



Atomic Multicast

→ Atomicity

- Message sent to a group arrives at all group members

- If it fails to arrive at any member, no member will process it

→ Problems

- Unreliable network

- Each message should be acknowledged

- Acknowledgements can be lost

- Message sender might die

How to achieve Atomicity?

→ General Idea

- Ensure that every recipient acknowledges receipt of the message
- Only then allow the application to process the message
- If we give up on a recipient then no recipient can process the received message

→ Easier said than done!

- What if a recipient dies after acknowledging the message?
 - Is it obligated to restart?
 - If it restarts, will it know to process the message?
- What if the sender (or coordinator) dies partway through the protocol?

Achieving Atomicity - An Example

Retry through network failures & system downtime

- Sender & receivers maintain a **persistent log**
- Each message has a unique ID so we can discard duplicates
- **Sender**
 - sends the message to all group members
 - Writes the message to log
 - Waits for acknowledgement from each group member
 - Writes the acknowledgement to log
 - If timeout on waiting for an acknowledgement, retransmit to group member
- **Receiver** logs received non-duplicate message to persistent log and sends an acknowledgement

NEVER GIVE UP! - Assume that dead senders or receivers will be rebooted and will restart where they left off

Reliable multicast

All non-faulty group members will receive a message

- Assume sender & recipients will remain alive
- Network may have glitches
 - Retransmit undelivered messages

Acknowledgements

- Send message to each group member
- Wait for acknowledgement from each group member
- Retransmit to non-responding members
- Subject to feedback implosion

Negative acknowledgements

- Use a sequence number on each message
- Receiver requests retransmission of a missed message
- More efficient but requires sender to buffer messages indefinitely

Acknowledgements

- ➡ Easiest thing is to wait for an ACK before sending the next message
 - But that incurs a round-trip delay
- ➡ Optimizing
 - Pipelining
 - Send multiple messages - receive ACKs asynchronously
 - Set timeout - retransmit message for missing ACKs
 - Cumulative ACKs
 - Wait a little while before sending an ACK
 - If you receive others, then send one ACK for everything
 - Piggybacked ACKs
 - Send an ACK along with a return message
- ➡ TCP does all of these ... but now we have to do this on each recipient

Message Ordering

➡ How to order messages?

- ➡ Send vs Delivery

- ➡ Global Time Ordering

- ➡ Total Ordering

- ➡ Causal Ordering


- ➡ Sync Ordering

- ➡ FIFO Ordering


- ➡ Unordered multicast

➡ Good / Bad Ordering


Good Ordering




Bad Ordering



Good Ordering



Bad Ordering



Send vs. Delivery of Messages

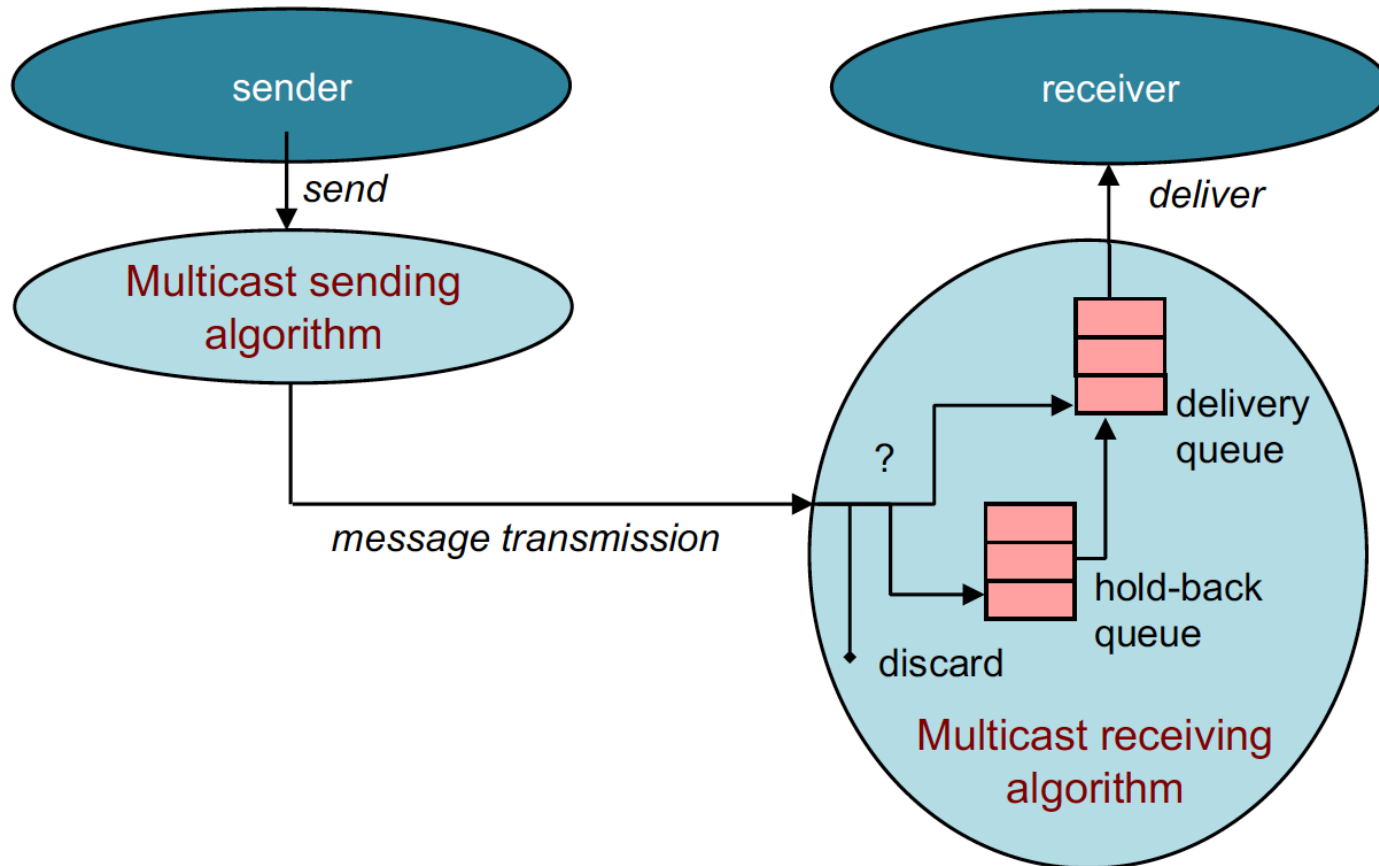
- ➡ Multicast receiver algorithm decides when to deliver a message to a process

A received message may be:

- ➡ delivered immediately (put on a delivery queue that the process reads)
- ➡ placed on a hold-back queue (because we need to wait for an earlier message)
- ➡ rejected/discarded (duplicate or earlier message that we no longer want)

An Illustration

➡ Sending, delivering and holding back



Global Time Ordering

- ➡ All messages arrive in exact order sent
- ➡ Assumes that two events never happen at exactly the same time!
 - ➡ Why Not? No global clocks ... right?
- ➡ Difficult (impossible) to achieve

Total Ordering

- ➡ Consistent ordering everywhere
- ➡ All messages arrive at all group members in the same order
 - ➡ They are sorted in the same order in the delivery queue

Two Conditions:

- ➡ If a process sends m before m' then any other process that delivers m' will have delivered m
- ➡ If a process delivers m' before m'' then every other process will have delivered m' before m''

Total Ordering - Implementation

➡ How to implement this?

➡ Attach unique totally sequenced message ID

➡ Receiver delivers a message to the application only if it has received all messages with a smaller ID

Causal Ordering


➡ Partial ordering

- ➡ Messages sequenced by Lamport or Vector timestamps

Condition:


- ➡ If $\text{multicast}(G, m) \rightarrow \text{multicast}(G, m')$
 - ➡ then every process that delivers m' will have m delivered already
- ➡ If message m' is causally dependent on the message m , then all processes must deliver m before m'

Causal vs Concurrent



Causal

m_1 is causally dependent on the receipt of m_0 .
Hence, m_1 must be delivered after m_0 has been delivered.



Concurrent

m_0 and m_1 have no causal relationship (they are concurrent).
Any process can deliver them in any order.

Causal Ordering - Implementation

How to implement CO?

- ➡ P_i receives a message from P_j
- ➡ Each process keeps a precedence vector (similar to vector timestamp)
- ➡ Vector is updated on multicast send and receive events
 - ➡ Each entry = number of the latest message from the corresponding group member that causally precedes the event

Causal Ordering - Algorithm

- ➡ When P_j **sends** a message, it increments its own entry and sends the vector
 - ➡ $V_j[j] = V_j[j] + 1$
 - ➡ Send V_j with the message
- ➡ When P_i **receives** a message from P_j
 - ➡ Check that the message arrived in FIFO order from P_j :
 $V_j[j] == V_i[j] + 1$?
 - ➡ Check that the message does not causally depend on something P_i has not seen
 $\forall k, k \neq j: V_j[k] \leq V_i[k]$?
 - ➡ If both conditions are satisfied, P_i will deliver the message
 - ➡ Otherwise, hold the message until the conditions are satisfied

Causal Ordering - Work out


➡ Implementation:

P_i receives a message from P_j


➡ Each process keeps a precedence vector (similar to vector timestamp)

➡ Vector is updated on multicast **send** and **receive** events

➡ Each entry = Number of the latest message from the corresponding group member that causally precedes the event message




Causal Ordering - Example




- ➡ P_2 receives message m_1 from P_1 with $V_1=(1,1,0)$
- ➡ Is this in FIFO order from P_1 ?
 - ➡ Compare current V on P_2 : $V_2=(0,0,0)$ with received V from P_1 , $V_1=(1,1,0)$
 - ➡ Yes: $V_2[1] = 0$, received $V_1[1] = 1 \Rightarrow$ sequential order
- ➡ Is $V_1[i] \leq V_2[i]$ for all other i ?
 - ➡ Compare the same vectors: $V_2=(0,0,0)$ vs. $V_1=(1,1,0)$
 - ➡ No. $V_1[0] > V_2[0]$ ($1 > 0$)
 - ➡ Therefore: hold back m_1 at P_2

Causal Ordering - Example (contd)



- ➡ P_2 receives message m_0 from P_0 with $V=(1,0,0)$
- ➡ (1) Is this in FIFO order from P_0 ?
 - ➡ Compare current V on P_2 : $V_2=(0,0,0)$ with received V from P_2 , $V_2=(1,0,0)$
 - ➡ Yes: $V_2[0] = 0$, received $V_1[0] = 1 \Rightarrow$ sequential
- ➡ (2) Is $V_0[i] \leq V_2[i]$ for all other i ?
 - ➡ Yes
- ➡ Deliver m_0
 - ➡ Now check hold-back queue. Can we deliver m_1 ?

Causal Ordering - Example (contd)



- ➡ Is the held-back message m_1 in FIFO order from P_0 ?
 - ➡ Compare current V on P_2 : $V_2=(1,0,0)$ with held-back V from P_0 , $V_1=(1,1,0)$
 - ➡ Yes: $V_2[1] = 0$, received $V_1[1] = 1 \Rightarrow$ sequential
- ➡ Is $V_0[i] \leq V_2[i]$ for all other i ?
 - ➡ Now yes. Element 0: $(1 \leq 1)$, element 2: $(0 \leq 0)$; Deliver m_1
- ➡ More efficient than total ordering:
 - ➡ No need for a global sequencer.
 - ➡ No need to send acknowledgements.

Sync Ordering

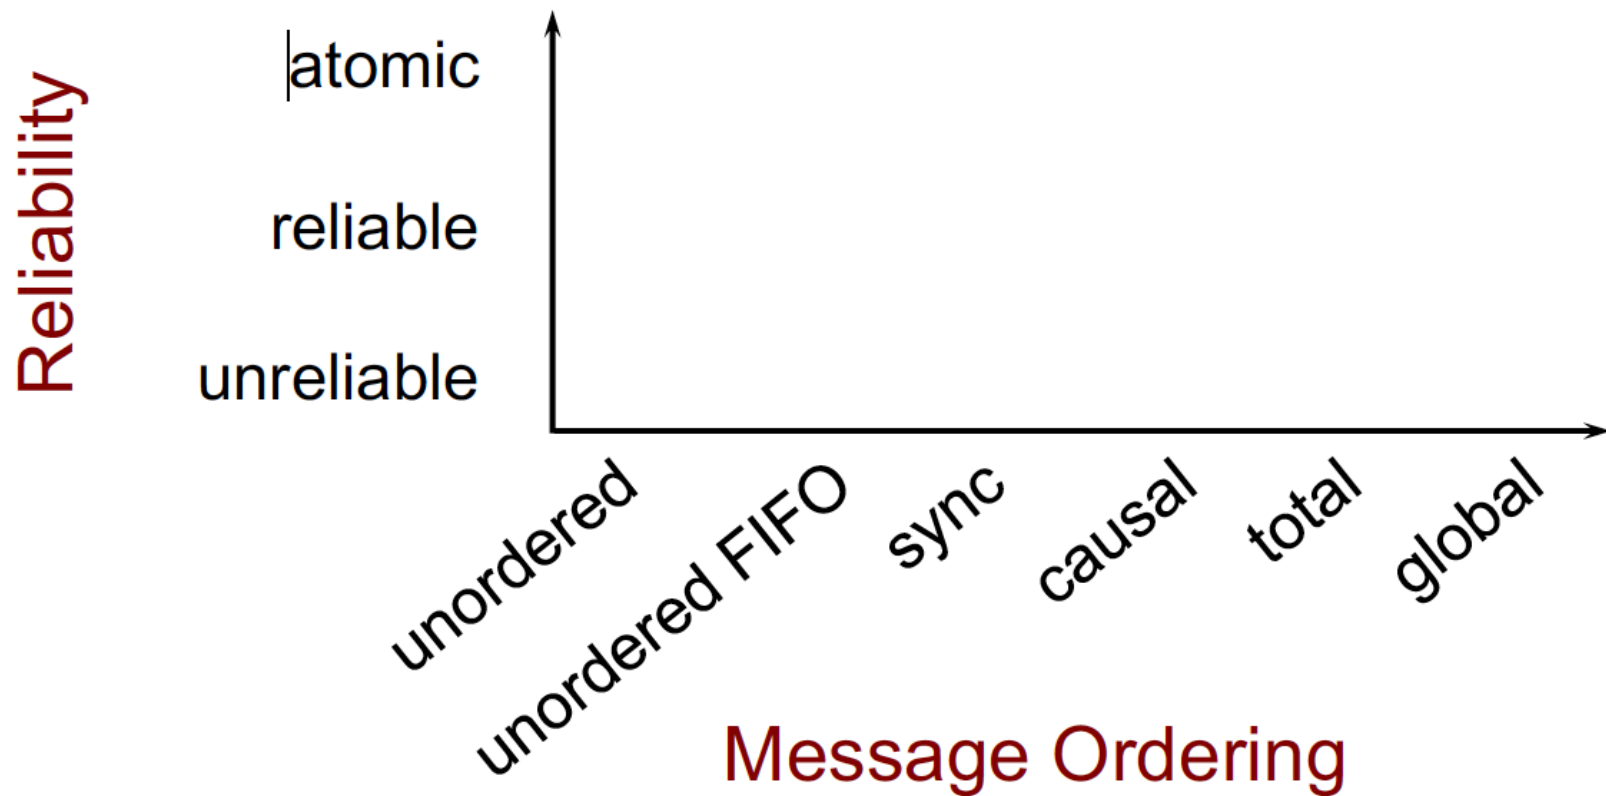
- ➡ Messages can arrive in any order
- ➡ Special message type
 - ➡ Synchronization primitive
 - ➡ Ensure all pending messages are delivered before any additional (post-sync) messages are accepted

Unordered multicast

- ➡ Messages can be delivered in different order to different members
- ➡ Order per-source does not matter

Multicast Considerations

➡ Follow this order !!



Summary

- Communication Models
- Message Ordering & Group Communications
- Design Issues
 - Process Failures
- Message Ordering
 - Good / Bad ordering
 - Various Types of Ordering of messages
- Group Communication
 - Causal ordering based approach
- Many more to come up ... stay tuned in !!

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>


OR

→ <http://rajendra.2power3.com>

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

Thanks ...



... Questions ???