

Distributed Mutual Exclusion Algorithms

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

About this topic

This course covers various concepts in **Mutual Exclusion in Distributed Systems**. We will also focus on different types of distributed mutual exclusion algorithms in distributed contexts and their analysis

What did you learn so far?

- Challenges in Message Passing systems
- Distributed Sorting
- Space-Time Diagram
- Partial Ordering / Causal Ordering
- Concurrent Events
- Local Clocks and Vector Clocks
- Distributed Snapshots
- Termination Detection
- Topology Abstraction and Overlays
- Leader Election Problem in Rings
- Message Ordering / Group Communications

Recent Topic ...

- Communication Models
- Design Issues
 - Process Failures
- Message Ordering
 - Good / Bad ordering
 - Various Types of Ordering of messages
- Group Communication
 - Causal ordering based approach
 - Many more to come up ... stay tuned in !!

Topics to focus on ...

- Leader Election in Distributed Systems
- Topology Abstraction and Overlays
- Message Ordering
- Group Communication
- Distributed Mutual Exclusion
- Deadlock Detection
- Check pointing and rollback recovery

For End Semester

Mutual Exclusion in Distributed Systems

Let us explore Mutex algorithms proposed for various interconnection networks

Distributed Mutual Exclusion

- **No Deadlocks** - no set of sites should be permanently blocked, waiting for messages from other sites in that set
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation.)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites

Distributed MutEx - An overview

Token-based solution: Processes share a special message known as a token

- ➔ Token holder has right to access shared resource
- ➔ Wait for/ask for (depending on algorithm) token; enter Critical Section (CS) when it is obtained, pass to another process on exit or hold until requested (depending on algorithm)
- ➔ If a process receives the token and doesn't need it, just pass it on

Distributed MutEx - A Few Issues

- Who can access the resource?
- When does a process to be privileged to access the resource?
- How long does a process access the resource?
Any finite duration?
- How long can a process wait to be privileged?
- Computation complexity of the solution

Types of Distributed MutEx

- Token-based distributed mutual exclusion algorithms
 - Suzuki - Kasami's Algorithm
- Non-token based distributed mutual exclusion algorithms
 - Lamport's Algorithm
 - Ricart-Agrawala's Algorithm

Token Based Methods

Advantages:

- Starvation can be avoided by efficient organization of the processes
- Deadlock is also avoidable

Disadvantage: Token Loss

- Must initiate a cooperative procedure to recreate the token
- Must ensure that only one token is created!

Non-Token Based Methods

- ➔ **Permission-based solutions:** a process that wishes to access a shared resource must first get permission from one or more other processes.
- ➔ Avoids the problems of token-based solutions, but is more complicated to implement

Token Ring Approach


- Processes are organized in a logical ring: P_i has a communication channel to $P_{(i+1) \bmod N}$

Operations:

- Only the process holding the token T can enter the CS
- To enter the critical section, wait passively for T
When in CS, hold on to T and don't release it
- To exit the CS, send T onto your neighbor
- If a process does not want to enter the CS when it receives T, it simply forwards T to the next neighbor

Token Rings - Illustration

➔ Request movements in an unidirectional ring network



Token Based Control in Rings

- Requests move in either Clockwise or Anticlockwise
- Proposed by Feuerstein et al. (1996)
- There are 3 steps:
 - P needs the resource
 - P has T: enter CS
 - P has no T: send the request to the next P
 - P receives a request
 - P has T: increase TC by 1 and send T to the next P
 - P has no T: send request to the next P
 - P receives Token
 - P has a pending request: enter CS and decrease TC by 1 and send T to next P if $TC > 0$
 - P has no pending request: send T to the next P

Refer to: Feuerstein et al., Efficient token-based control in rings, Information Processing Letters, 66 (4) (1998) pp. 175-180

Token Rings - Features

- Safety & Liveness are guaranteed
- Ordering is not guaranteed
- Bandwidth: 1 message per exit
- Client delay: 0 to N message transmissions
- Synchronization delay between one process's exit from the CS and the next process's entry is between 1 and N-1 message transmissions

Non-Token Based Algorithms

→ Notations:

→ P_i : i^{th} Process

→ R_i : Request set, containing IDs of all P_i s from which permission must be received before accessing CS

→ Non-token based approaches use time stamps to order requests for CS

→ Smaller time stamps get priority over larger ones

→ Lamport's Algorithm

→ $R_i = \{P_1, P_2, \dots, P_n\}$, i.e., all processes.

→ Request queue: maintained at each P_i ordered by time stamps.

→ **Assumption:** message delivered in FIFO

Lamport's Algorithm

→ Requesting CS:

- Send REQUEST(ts_i, i) where (ts_i, i) - Request time stamp; Place REQUEST in $request_queue_i$
- On receiving the message; P_j sends time-stamped REPLY message to P_i ; P_i 's request placed in $request_queue_j$

→ Executing CS:

- P_i has received a message with time stamp larger than (ts_i, i) from all other sites
- P_i 's request is the top most one in $request_queue_i$

→ Releasing CS:

- Exiting CS: send a time stamped RELEASE message to all sites in its request set
- Receiving RELEASE message: P_j removes P_i 's request from its queue

Notable Points

- Purpose of REPLY messages from i to j is to ensure that j knows of all requests of i prior to sending the REPLY (possibly any request of i with timestamp lower than j 's request)
- Requires FIFO channels
- $3(n - 1)$ messages per critical section invocation
- Synchronization delay = max msg transmission time
- Requests are granted in order of increasing timestamps

Performance Improvements

- $3(n-1)$ messages per Critical Section invocation
 - $(n - 1)$ REQUEST messages
 - $(n - 1)$ REPLY messages
 - $(n - 1)$ RELEASE messages
- Synchronization delay: T
- **Optimization:**
 - Suppress reply messages: For example, P_j receives a REQUEST message from P_i after sending its own REQUEST message with time stamp higher than that of P_i 's then Do NOT send a REPLY message
 - Messages reduced to between $2(n-1)$ and $3(n-1)$

Ricart & Agrawala's Algorithm

- A time-stamp based approach
- Originally proposed by Lamport using logical clocks
- Modified by Ricart & Agrawala

Ricart & Agrawala's Algorithm

Main Idea:

- Process j need not send a REPLY to Process i if j has a request with timestamp lower than the request of i (since i cannot enter before j here)
- Does not require FIFO
- $2(n - 1)$ messages per critical section invocation
- Synchronization delay = maximum message transmission time
- Requests granted in order of increasing timestamps

Ricart & Agrawala (contd)

- Processes need entry to critical section multicast a request, and can enter it only when all other processes have replied positively
- Messages requesting entry are of the form $\langle T, P_i \rangle$
 - T - sender's timestamp (Lamport clock)
 - P_i the sender's identity

Ricart & Agrawala - Algorithm

To enter the Critical Section (CS):

- Set state = wanted
- multicast "request" to all processes (including timestamp)
- wait until all processes send back "reply"
- change state to held and enter the CS

On receipt of a request $\langle T_j, P_j \rangle$ at P_i :

- if (state == held) or (state == wanted & $(T_i, P_i) < (T_j, P_j)$) then enqueue the request
- else "reply" to P_j

On exiting the CS:

- change state to release and "reply" to all queued requests

Ricart & Agrawala - Simplified

To request Critical Section:

- send timestamped REQUEST message (ts_i, i)

On receiving request (ts_i, i) at j :

- if j is neither requesting nor executing critical section then send REPLY to i
- if j is requesting and i 's request timestamp is smaller than j 's request timestamp then
 - enqueue the request; Otherwise, defer the request

To enter Critical Section:

- Process i enters critical section on receiving REPLY messages from all processes

To release Critical Section:

- send REPLY to all deferred requests

Summary

- Mutual Exclusion Problem
- Basics of MutEx algorithms
- Various Types of MutEx algorithms
 - Token-based
 - Token rings
 - Non-Token based algorithm
 - Lamport's Algorithm
 - Ricart - Agrawala's Algorithm
- Performance Metrics
 - Many more to come up ... stay tuned in !!

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>


OR

→ <http://rajendra.2power3.com>

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

Thanks ...



... Questions ???