

Distributed Mutual Exclusion Algorithms

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

About this topic

This course covers various concepts in **Mutual Exclusion in Distributed Systems**. We will also focus on different types of distributed mutual exclusion algorithms in distributed contexts and their analysis

What did you learn so far?

- Challenges in Message Passing systems
- Distributed Sorting
- Space-Time Diagram
- Partial Ordering / Causal Ordering
- Concurrent Events
- Local Clocks and Vector Clocks
- Distributed Snapshots
- Termination Detection
- Topology Abstraction and Overlays
- Leader Election Problem in Rings
- Message Ordering / Group Communications

Topics to focus on ...

→ Distributed Mutual Exclusion

- Deadlock Detection
- Check pointing and rollback recovery
- Self-Stabilization
- Distributed Consensus
- Reasoning with Knowledge
- Peer – to – peer computing and Overlays
- Authentication in Distributed Systems

Mutual Exclusion in Distributed Systems

Let us explore mutex algorithms proposed for various interconnection networks

Why do we need Mutex?

→ Mutual Exclusion

→ Operating systems: Semaphores

→ In a single machine, you could use semaphores to implement mutual exclusion

→ How to implement semaphores?

→ Inhibit interrupts

→ Use clever instructions (e.g. test-and-set)

→ On a multiprocessor shared memory machine, only the latter works

Characteristics

- Processes communicate only through messages - no shared memory or no global clocks
- Processes must expect unpredictable message delays
- Processes coordinate access to shared resources (printer, file, etc.) that should be used in a mutually exclusive manner

Race Conditions

- Consider Online systems – For example, Airline reservation systems maintain records of available seats
- Suppose two people buy the same seat, because each checks and finds the seat available, then each buys the seat
- Overlapped accesses generate different results than serial accesses
 - race condition

Distributed Mutual Exclusion

→ Needs

- Only one process should be in critical section at any point of time
- What about resources?

Distributed Mutual Exclusion

- **No Deadlocks** - no set of sites should be permanently blocked, waiting for messages from other sites in that set
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation.)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites

Distributed MutEx - An overview

Token-based solution: Processes share a special message known as a token

- ➔ Token holder has right to access shared resource
- ➔ Wait for/ask for (depending on algorithm) token; enter Critical Section (CS) when it is obtained, pass to another process on exit or hold until requested (depending on algorithm)
- ➔ If a process receives the token and doesn't need it, just pass it on

Distributed MutEx - A Few Issues

- Who can access the resource?
- When does a process to be privileged to access the resource?
- How long does a process access the resource? Any finite duration?
- How long can a process wait to be privileged?
- Computation complexity of the solution

Types of Distributed MutEx

- Token-based distributed mutual exclusion algorithms
 - Suzuki - Kasami's Algorithm
- Non-token based distributed mutual exclusion algorithms
 - Lamport's Algorithm
 - Ricart-Agartala's Algorithm

Token Based Methods

Advantages:

- Starvation can be avoided by efficient organization of the processes
- Deadlock is also avoidable

Disadvantage: Token Loss

- Must initiate a cooperative procedure to recreate the token
- Must ensure that only one token is created!

Non-Token Based Methods

- ➔ **Permission-based solutions:** a process that wishes to access a shared resource must first get permission from one or more other processes.
- ➔ Avoids the problems of token-based solutions, but is more complicated to implement

Performance Analysis

- Guarantees mutual exclusion
- No starvation: Only if requests served in order
- No deadlock
- Fault tolerant?
 - Single point of failure
 - Blocking requests mean client processes have difficulty distinguishing crashed coordinator from long wait
 - Bottlenecks
- The solution is simple and ease

Quorum Based algorithms

Why Quorum based algorithm?

- Lamports and Ricard-Agrawala' algorithm requires permission from all processes to enter into the critical section.

Modifications:

- Is it necessary to obtain permission from all processes before entering into the CS?
- How to reduce the message exchanges and increase the performance of Mutex algorithm?

Quorum Based algorithms

What is a Quorum?


- There are n requesting processes in a distributed system and any process may request for CS.
- Can we form such a subset of processes who request for Critical Section? YES !!
- Such a set is said to be a Request Set or Quorum
- In fact, we will have a separate Request set for each process P_i

Quorum - Definition

- A quorum system is a collection of subsets of processes, called quorums, such that each pair of quorums have a non-empty intersection
- How do we formally define a quorum of processes in a distributed system?
- Let us look at some examples

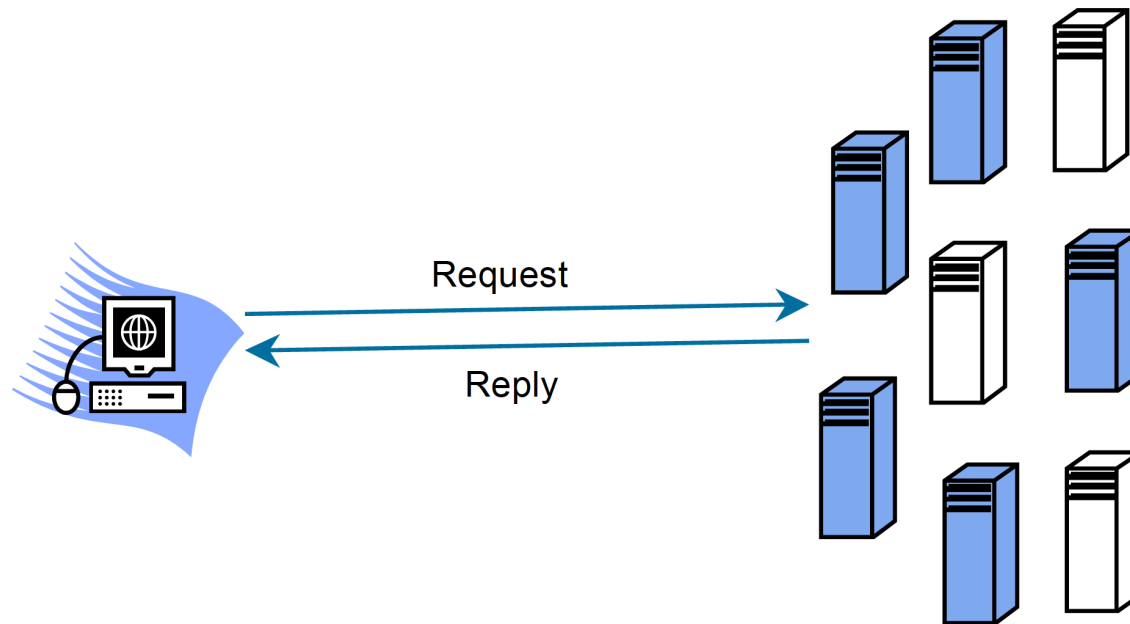
Quorum - Why?

- Process may not respond or may go down (any kind of failure)
- The requesting process can not get REPLY from all remaining processes
- It would infinitely wait for CS !!



Quorum - Why?

→ Can the requesting process get permission from a quorum of processes to enter into CS?



Quorum - Definition

More Formally,

→ Given a set of processes

$$P = \{P_1, P_2, \dots, P_n\}$$

→ A quorum system $Q \subseteq 2^P$ is a set of subsets of P such that

for all Q_1, Q_2 in Q : $Q_1 \cap Q_2 \neq \text{empty}$

→ Each Q_i in Q is called a quorum

Maekawa's Algorithm

- Permission obtained from only a subset of other processes, called the Request Set (or Quorum)
- Separate Request Set R_i , for each process i

Maekawa's Algorithm

Requirements

- For all $i, j: R_i \cap R_j \neq \Phi$
- For all $i: i \in R_i$
- For all $i: |R_i| = K$, for some K
- Any node i is contained in exactly D Request Sets, for some Request set D
- $K = D = \text{sqrt}(N)$ for Maekawa's algorithm

Maekawa's Algorithm - Steps

To Request Critical Section:

→ P_i sends REQUEST message to all process in R_i

On receiving a REQUEST message:

→ Send a REPLY message if no REPLY message has been sent since the last RELEASE message is received.

→ Update status to indicate that a REPLY has been sent.

→ Otherwise, queue up the REQUEST

To enter critical section:

→ P_i enters critical section after receiving REPLY from all nodes in R_i

Maekawa's Algorithm - Steps (contd)

To release critical section:

- Send RELEASE message to all nodes in R_i
- On receiving a RELEASE message, send REPLY to next node in queue and delete the node from the queue.
- If queue is empty, update status to indicate no REPLY message has been sent

Computation Complexity


- Message Complexity: $3 * \sqrt{N}$
- Synchronization delay
 - $2 * (\text{max message transmission time})$
- Major problem: DEADLOCK possible
- Need three more types of messages (FAILED, INQUIRE, YIELD) to handle deadlock.
 - Message complexity can be $5 * \sqrt{N}$
- Important Issue:
 - How to build the request sets?

Token Ring Approach

- Single token circulates, enter CS when token is present
- Mutual exclusion obvious
- Algorithms differ in how to find and get the token
- Uses sequence numbers rather than timestamps to differentiate between old and current requests

Token Rings - Illustration

➔ Request movements in an unidirectional ring network



Suzuki - Kasami's Algorithm

- Broadcast a request for the token
- Process with the token sends it to the requestor if it does not need it
- **Issues:**
 - Current versus outdated requests
 - Determining sites with pending requests
 - Deciding which site to give the token to

Data Structures

The token:

- Queue (FIFO) Q of requesting processes
- $LN[1..n]$: sequence number of request that j executed most recently

The request message:

- $REQUEST(i, k)$: request message from node i for its k^{th} critical section execution

Other data structures:

- $RN_i[1..n]$ for each node i , where $RN_i[j]$ is the largest sequence number received so far by i in a $REQUEST$ message from j

Suzuki-Kasami's algorithm

To request critical section:

- If i does not have token, increment $RN_i[i]$ and send $REQUEST(i, RN_i[i])$ to all nodes
- If i has token already, enter critical section if the token is idle (no pending requests), else follow rule to release critical section

On receiving $REQUEST(i, s_n)$ at j :

- Set $RN_j[i] = \max(RN_j[i], s_n)$
- If j has the token and the token is idle then
 - send it to i if $RN_j[i] = LN[i] + 1$
 - If token is not idle, follow rule to release critical section

Suzuki-Kasami's algorithm

To enter critical section:

→ Enter CS if token is present

To release critical section:

→ Set $LN[i] = RN_i[i]$

→ For every node j which is not in Q (in token),
add node j to Q if $RN_i[j] = LN[j] + 1$

→ If Q is non empty after the above, delete first
node from Q and send the token to that node

Complexity

→ No. of messages:

- 0 if node holds the token already,
- n otherwise

→ Synchronization delay:

- 0 (node has the token) or
- max. message delay (token is elsewhere)

→ No starvation

Raymond's Algorithm

- Forms a directed tree (logical) with the token token-holder as root
- Each node has variable "Holder" that points to its parent on the path to the root.
 - Root's Holder variable points to itself
- Each node P_i has a FIFO request queue Q_i

Raymond's Algorithm

- To request critical section:
 - Send REQUEST to parent on the tree, provided i does not hold the token currently and Q_i is empty. Then place request in Q_i
- When a non-root node j receives a request from k
 - place request in Q_j
 - send REQUEST to parent if no previous REQUEST sent

Raymond's Algorithm (contd)

When the root receives a REQUEST:

- send the token to the requesting node
- set Holder variable to point to that node

When a node receives the token:

- delete first entry from the queue
- send token to that node
- set Holder variable to point to that node
- if queue is non non-empty, send a REQUEST message to the parent (node pointed at by Holder variable)

Raymond's Algorithm (contd)

→ To execute critical section:

- enter if token is received and own entry is at the top of the queue; delete the entry from the queue

→ To release critical section:

- if queue is non non-empty, delete first entry from the queue, send token to that node and make Holder variable point to that node
- If queue is still non non-empty, send a REQUEST message to the parent (node pointed at by Holder variable)

Features of Raymond's Algo

- Average message complexity:
 - $O(\log n)$
- Sync. Delay
 - $(T \log n)/2$, where T = max. message delay

Summary

→ Mutual Exclusion

→ Various Types of MutEx algorithms

→ Non-Token based algorithm

→ Quorum based algorithm

→ Token based algorithm

→ Suzuki - Kasami's Algorithm

→ Raymond's Tree based algorithm

→ Performance Metrics

→ Stay tuned ... More to come up ... !!

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>


OR

→ <http://rajendra.2power3.com>

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

Thanks ...



... Questions ???