

# Deadlocks Detection

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

# About this topic

This course covers various concepts in **Mutual Exclusion in Distributed Systems**. We will also focus on different types of distributed mutual exclusion algorithms in distributed contexts and their analysis

# What did you learn so far?

- Challenges in Message Passing systems
- Distributed Sorting
- Space-Time Diagram
- Partial Ordering / Causal Ordering
- Concurrent Events
- Local Clocks and Vector Clocks
- Distributed Snapshots
- Termination Detection
- Topology Abstraction and Overlays
- Leader Election Problem in Rings
- Message Ordering / Group Communications
- Distributed Mutual Exclusion Algorithms

# Topics to focus on ...

- Distributed Mutual Exclusion
- **Deadlock Detection**
- Check pointing and rollback recovery
- Self-Stabilization
- Distributed Consensus
- Reasoning with Knowledge
- Peer - to - peer computing and Overlays
- Authentication in Distributed Systems

# Deadlocks



Let us explore deadlock detection, prevention and avoidance algorithms in distributed systems

# Distributed Mutual Exclusion (recap)

- **No Deadlocks** - No processes should be permanently blocked, waiting for messages (Resources) from other sites
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites


# Deadlock - A Simple Example

→ Vehicular Traffic at a signal



# Deadlock - Another Example

→ Vehicular Traffic - Another Scenario





# Deadlock - Illustrated

→ Vehicular Traffic - A real-time scenario



# Dining Philosophers' Problem

- Each philosopher must alternately think and eat
- A philosopher can only eat when they have both left and right forks
- **Problem:** How to design a discipline of behavior (a concurrent algorithm) such that no philosopher will starve?

→ Suggest a Simple Solution ??



# Dining Philosophers' Problem

- **Soln - 1:** Forks will be numbered 1 through 5 and each philosopher will always pick up the lower-numbered fork first, and then the higher-numbered fork
- **Soln - 2:** Use Arbitrator (waiter) to grant permission to pick up both forks

→ **Deadlock-Free Solutions !!**



# Deadlocks in Distributed Systems

## Definition

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set
- No process can progress in the system
- Competing processes may **WAIT** indefinitely for resources
- How do we manage resources among the competing tasks efficiently?

# Deadlocks - A few more examples

## Tape Drives

- Assume that a system has two Tape Drives
- There are two processes  $P_1$  and  $P_2$  each hold one drive
- Now each process needs access to another tape drive
- $P_1$  does not get access to the resource held by  $P_2$  and vice versa.
- This implies **DEADLOCK** ... neither  $P_1$  nor  $P_2$  succeeds in its attempt

# Deadlocks - A few more examples

## Semaphores

→ Semaphores A and B

$P_1$   
wait (A)

$P_2$   
wait(B)

OR

wait (B)

wait(A)

→ This implies **DEADLOCK** ... neither  $P_1$  nor  $P_2$  succeeds in its attempt

# Deadlock - Characterization

- **Mutual exclusion** - only one process at a time can use a resource
- **Hold and wait** - a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption** - a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait** - there exists a set  $\{P_0, P_1, \dots, P_{n-1}\}$  of waiting processes such that  $P_i$  is waiting for a resource that is held by  $P_{j \pmod n}$  where  $n$  is the total number of resources

# System Model

- Resource types:  $R_1, R_2, \dots, R_m$ 
  - CPU cycles, memory space, I/O devices
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - REQUEST
  - USE (Critical Section)
  - RELEASE
- Recall - Distributed Exclusion Algorithms



# Resource Allocation Graph (RAG)

→ A set of vertices  $V$  and a set of edges  $E$

→  $V$  is partitioned into two types:

→ Set consisting of all processes

$$P = \{P_1, P_2, \dots, P_n\}$$

→ Set consisting of all resource types


$$R = \{R_1, R_2, \dots, R_m\}$$

→ **request edge** - directed edge  $P_i \rightarrow R_j$


→ **assignment edge** - directed edge  $R_j \rightarrow P_i$

# Resource Allocation Graph (contd)


→ Process




→ Resource type with 4 instances



→  $P_i$  requests an instance of  $R_j$



→  $P_i$  is holding an instance of  $R_j$



# RAG - An example

→ Look at this graph

→ Resources:

→  $R_1$  - 1 unit


→  $R_2$  - 2 units

→  $R_3$  - 1 unit

→  $R_4$  - 3 units


→ Requests:

→  $P_1, P_2, P_3$




# RAG with a Deadlock

- Look at this graph
  - $P_1$  needs  $R_1$  which in turn used by  $P_2$  and  $P_2$  is requesting  $R_3$  which is currently being accessed by  $P_3$  and  $P_3$  needs  $R_2$  which is being locked by  $P_1$  and  $P_2$
- This implies **Deadlock**



# RAG with a cycle but NO Deadlock

- Look at this graph
  - $P_2$  and  $P_4$  may release the resource  $R_3$  in finite time as they do not depend on other competing processes
  - There exists a cycle but may not be a deadlock !!



# Basic Facts

- If graph contains no cycles → no deadlock
- If graph contains a cycle →
  - if **only one instance** per resource type, then deadlock
  - if **several instances** per resource type, possibility of deadlock

# How to handle Deadlocks?

- Ensure that the system will never enter a deadlock state
  - **Deadlock Prevention** - Stop before it happens!
  - **Deadlock Avoidance** - Precautions !!
  - **Deadlock Detection** - How to overcome?
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX

# Deadlock Prevention

## 4 Conditions to occur Deadlocks:

- **Mutual Exclusion** - Exclusive access - when a process accesses a resource, it is granted exclusive use of that resource
- **Hold and wait** - a process is allowed to hold onto some resources while waiting for other resources
- **No preemption** - a process cannot preempt or take away the resources held by another process
- **Cyclical wait** - There is a circular chain of waiting processes, each waiting for a resource held by the next process in the chain



# Deadlock Avoidance

Requires that the system has some additional a priori information available

- Simplest and most useful model requires that each process declare the maximum number of resources of each type that it may need
- Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition

# Safe State


- When a process requests an available resource, system must decide whether the allocation immediate leaves the system in a **Safe State**?
- System is in **Safe State** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL processes such that for each  $P_i$ , the resources that  $P_i$  can still request, can be satisfied by available resources + resources held by all  $P_j, j < i$
- If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
- When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
- When  $P_i$  terminates,  $P_{i+1}$  can get resources and so on

# Basic Facts

- If a system is in safe state
  - no deadlocks
- If a system is in unsafe state
  - possibility of deadlock
- Avoidance
  - ensure that a system will never enter an unsafe state

# Safe / Unsafe / Deadlock State

→ Illustration of safe, unsafe and deadlock state



# Deadlock Avoidance Algorithms

# Avoidance Algorithms

- Single instance of a resource type
  - Use a resource-allocation graph
- Multiple instances of a resource type
  - Use the Banker's algorithm



# Resource-Allocation Graph Scheme

- Claim edge  $P_i \rightarrow R_j$  indicate:
  - (i) process  $P_j$  may request resource  $R_j$
  - (ii) represented by a dashed line
- Claim edge converts to Request edge when a process requests for a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge **reconverts** to a claim edge
- Resources must be claimed a priori in the system

# Resource-Allocation Graph

→ An Example

Unsafe state





# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# Banker's Algorithm

- Multiple instances of Resources
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

# Data Structures - Banker's Algorithm

Let  $n$  = number of processes;  
 $m$  = number of resources types;

- **Available:** Vector of length  $m$ . If  $\text{available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $\text{Max}[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

# Safety Algorithm

- (1) Let **Work** and **Finish** be vectors of length  $m$  and  $n$  respectively. Initialize:  $\text{Work} = \text{Available}$  and  $\text{Finish}[i] = \text{false}$  for  $i = 0, 1, \dots, n-1$
- (2) Find an  $i$  such that both:
  - (a)  $\text{Finish}[i] = \text{false}$ ; (b)  $\text{Need}_i \leq \text{Work}$ .If no such  $i$  exists, go to step 4
- (3)  $\text{Work} = \text{Work} + \text{Allocation}_i$   
 $\text{Finish}[i] = \text{true}$ ; go to step 2
- (4) If  $\text{Finish}[i] == \text{true}$  for all  $i$ , then the system is in a safe state

# Resource-Request Algo for Process $P_i$

$Request_i$  = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $Request_i \leq Need_i$  then go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $Request_i \leq Available$  then go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe → the resources are allocated to  $P_i$
- If unsafe →  $P_i$  must wait, and the old resource-allocation state is restored

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types:
  - A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time  $T_0$ :

	Allocation	Max	Available
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

# Example (contd)

→ Need matrix is defined to be as follows:

$$\text{Need} = \text{Max} - \text{Allocation}$$

	Need		
	A	B	C
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	1	1
P4	4	3	1

→ The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria

# Example: $P_1$ Request (1,0,2)

→ Check that Request  $\leq$  Available

(that is, (1,0,2) (3,3,2) - Is true?

	Allocation	Need	Available
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

→ Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement

→ Can request for (3,3,0) by  $P_4$  be granted?

→ Can request for (0,2,0) by  $P_0$  be granted?




# Deadlock Detection

- Allow system to enter deadlock state
- Detection Algorithm
- Recovery Scheme


# Single Instance of Each Resource Type

- Maintain wait-for graph
- Nodes are processes
- $P_i \rightarrow P_j$  if  $P_i$  is waiting for Resources from  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph

# Resource-Allocation / Wait-for Graph



Resource allocation Graph



Wait For Graph (WFG)

# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates the current request of each process.

If Request  $[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

- (1) Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively Initialize:
  - (a) **Work** = Available
  - (b) For  $i = 1, 2, \dots, n$ ,  
    **Finish**[ $i$ ] = false
- (2) Find an index  $i$  such that both:
  - (a) **Finish**[ $i$ ] == false
  - (b) **Request** $_i \leq$  **Work**If no such  $i$  exists, go to step 4

# Detection Algorithm (contd)

- (3)  $Work = Work + Allocation_i$   
 $Finish[i] = true$ ; go to step 2
- (4) If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked

Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state

# Detection Algorithm - Example

→ Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)

→ Snapshot at time  $T_0$ :

	Allocation	Request	Available
	A B C	A B C	A B C
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

→ Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  
 $\text{Finish}[i] = \text{true}$  for all  $i$

# Detection Algo - Example (contd)

- $P_2$  requests an additional instance of type C

	Request			
	A	B	C	
$P_0$	0	0	0	
$P_1$	2	0	2	
$P_2$	0	0	1	
$P_3$	1	0	0	
$P_4$	0	0	2	

- State of system?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes' requests
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$



# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
- How often a deadlock is likely to occur?
- How many processes will need to be rolled back?
  - one for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

# Recovery from Deadlock: Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?

# Recovery from Deadlock: resource Preemption

- Selecting a victim – minimize cost
- Rollback – return to some safe state, restart process for that state
- Starvation – same process may always be picked as victim, include number of rollback in cost factor

# Resource Links

- Distributed Deadlock Detection
  - [http://www.cse.scu.edu/~jholliday/dd\\_9\\_16.htm](http://www.cse.scu.edu/~jholliday/dd_9_16.htm)
- Coffman et. al., System Deadlocks, ACM Computing Surveys. 3 (2) (1971): 67-78.  
DOI:10.1145/356586.356588
- Havender, James W., Avoiding deadlock in multitasking systems, IBM Systems Journal. 7 (2) (1968): 74.  
DOI:10.1147/sj.72.0074
- Knapp, Edgar, Deadlock detection in distributed databases, ACM Computing Surveys, 19 (4) (1987): 303-328. DOI:10.1145/45075.46163. ISSN 0360-0300

# Summary

## → Deadlocks

- Deadlock Prevention

- Deadlock Avoidance

- Deadlock Detection

  - Resource Allocation Graphs

  - Banker's Algorithm

- Recovery from Deadlocks

- Performance Metrics

  - Stay tuned ... More to come up ... !!

# How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>


OR

→ <http://rajendra.2power3.com>

# Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
  - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

# Thanks ...



## ... Questions ???